

# Artificial Intelligence

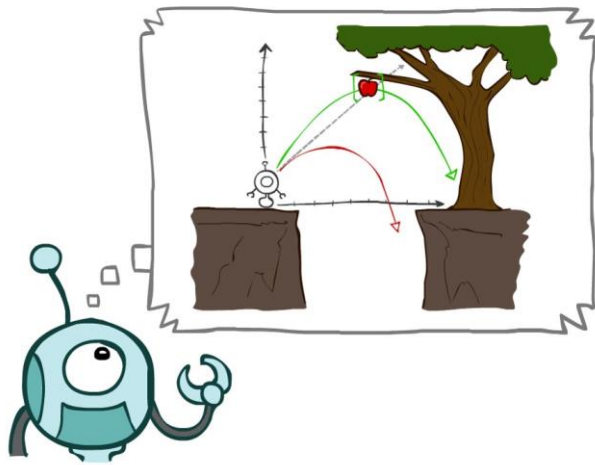
## CSE 4617

Ahnaf Munir

Assistant Professor

Islamic University of Technology

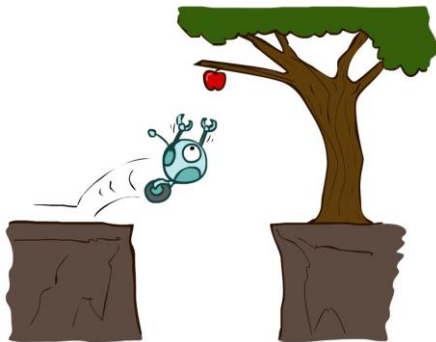
# Agents That Plan



# Reflex Agents

## ■ Properties

- Choose action based on current percept (and maybe memory)
- May have memory or model of the world's current state
- Do not consider future consequences
- Consider how the world **IS**
- Can be useful when quick decision is a must

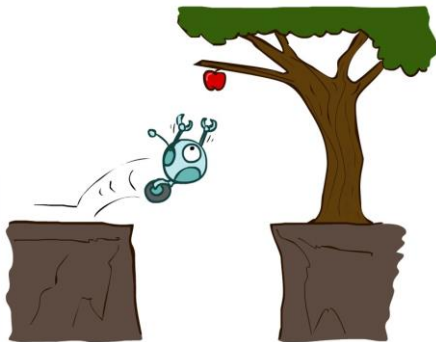


# Reflex Agents

## ■ Properties

- Choose action based on current percept (and maybe memory)
- May have memory or model of the world's current state
- Do not consider future consequences
- Consider how the world **IS**
- Can be useful when quick decision is a must

## ■ Can a reflex agent be rational?



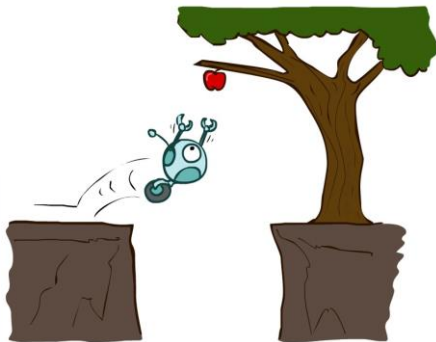
# Reflex Agents

## ■ Properties

- Choose action based on current percept (and maybe memory)
- May have memory or model of the world's current state
- Do not consider future consequences
- Consider how the world **IS**
- Can be useful when quick decision is a must

## ■ Can a reflex agent be rational?

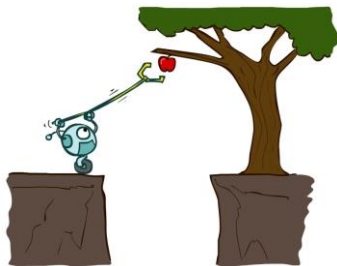
- Remember we only consider the outcome, not the process
- Only if quick decision is optimal



# Planning Agents

## ■ Properties

- Ask "what if"
- Decision based on (hypothesized) consequences of actions
- Must have a model of how the world evolves in response to actions
- Must formulate a goal (test)
- Consider how the world **WOULD BE**

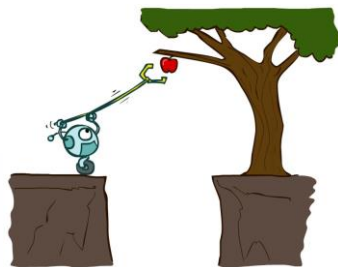


# Planning Agents

## ■ Properties

- Ask "what if"
- Decision based on (hypothesized) consequences of actions
- Must have a model of how the world evolves in response to actions
- Must formulate a goal (test)
- Consider how the world **WOULD BE**

- **Optimal** → Achieve goal in minimum cost
- **Complete** → When there exists a solution, find it



# Planning Agents


## ■ Properties

- Ask "what if"
- Decision based on (hypothesized) consequences of actions
- Must have a model of how the world evolves in response to actions
- Must formulate a goal (test)
- Consider how the world **WOULD BE**

■ **Optimal** → Achieve goal in minimum cost

**Complete** → When there exists a solution, find it

## ■ Planning vs Replanning





# Search Problems



# Search Problems

- State space  $\rightarrow$  Set of possible scenarios

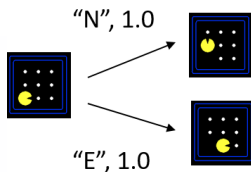


# Search Problems

- State space  $\rightarrow$  Set of possible scenarios



- A successor function (with actions, cost, etc.)  $\rightarrow$  Consequent states for given state

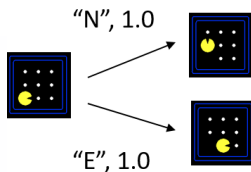


# Search Problems

- State space  $\rightarrow$  Set of possible scenarios



- A successor function (with actions, cost, etc.)  $\rightarrow$  Consequent states for given state



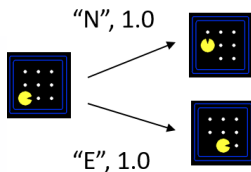
- A start state and goal test

# Search Problems

- State space → Set of possible scenarios



- A successor function (with actions, cost, etc.) → Consequent states for given state




- A start state and goal test
- Solution → A sequence of actions (a plan) which transforms the start state to goal state

# Search Problems Are Models



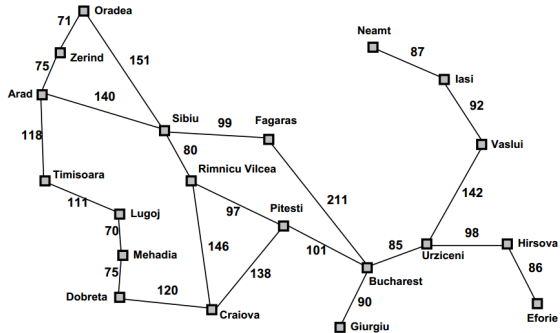
# Example: Travelling in Romania

■ State space?



# Example: Travelling in Romania

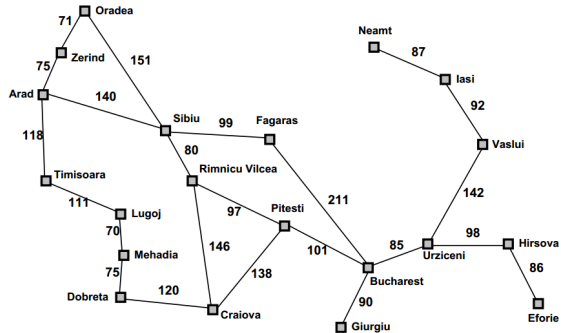
- State space?
- Cities



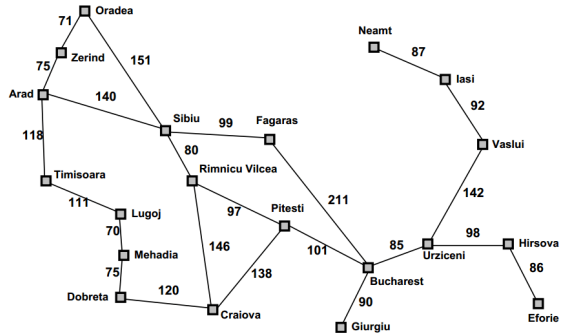


# Example: Travelling in Romania

- State space?
  - Cities
- Successor function?

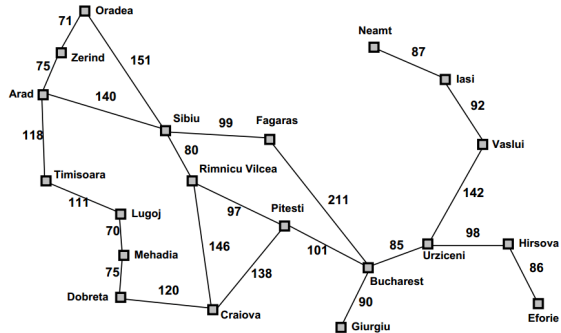


# Example: Travelling in Romania



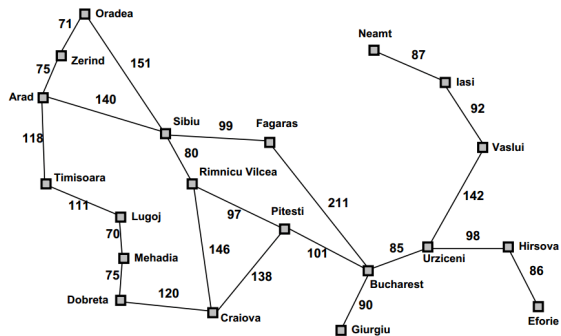
- State space?
  - Cities
- Successor function?
  - Roads: Go to adjacent city with cost = distance

# Example: Travelling in Romania



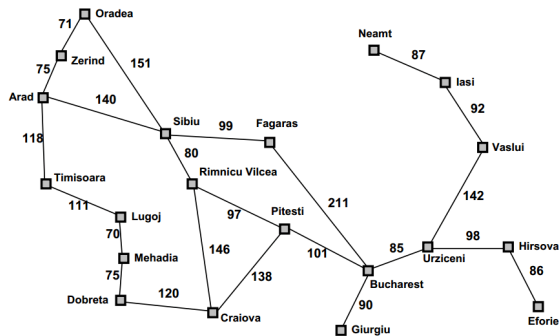
- State space?
  - Cities
- Successor function?
  - Roads: Go to adjacent city with cost = distance
- Start state?

# Example: Travelling in Romania




- State space?
  - Cities
- Successor function?
  - Roads: Go to adjacent city with cost = distance
- Start state?
  - Arad

# Example: Travelling in Romania



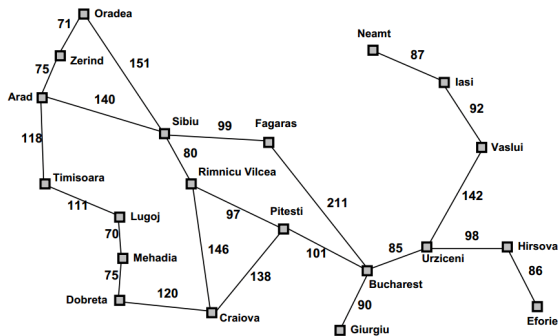
- State space?
  - Cities
- Successor function?
  - Roads: Go to adjacent city with cost = distance
- Start state?
  - Arad
- Goal test?

# Example: Travelling in Romania




- State space?
  - Cities
- Successor function?
  - Roads: Go to adjacent city with cost = distance
- Start state?
  - Arad
- Goal test?
  - Is state == Bucharest?

# Example: Travelling in Romania



- State space?
  - Cities
- Successor function?
  - Roads: Go to adjacent city with cost = distance
- Start state?
  - Arad
- Goal test?
  - Is state == Bucharest?
- Solution?

# Example: Travelling in Romania



## ■ State space?

- Cities

## ■ Successor function?

- Roads: Go to adjacent city with cost = distance

## ■ Start state?

- Arad

## ■ Goal test?

- Is state == Bucharest?


## ■ Solution?

- Path from Arad to Bucharest



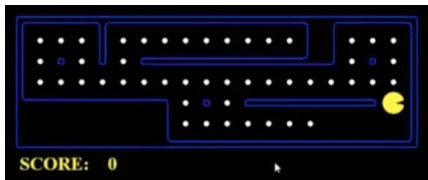
# What's in a State Space?

The *world state* includes every last detail of the environment



# What's in a State Space?

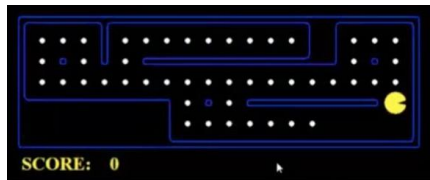
The *world state* includes every last detail of the environment



The *search state* keeps only the details needed for planning

# What's in a State Space?

The *world state* includes every last detail of the environment



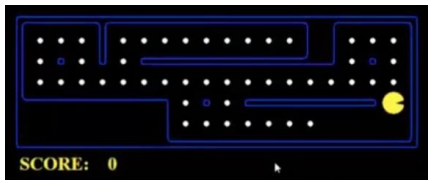
The *search state* keeps only the details needed for planning

## ■ Problem: Pathing

- States:

# What's in a State Space?

The *world state* includes every last detail of the environment

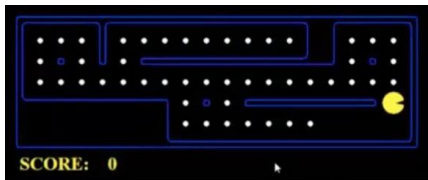


The *search state* keeps only the details needed for planning

- Problem: Pathing
  - States:  $(x, y)$  location
  - Actions:

# What's in a State Space?

The *world state* includes every last detail of the environment



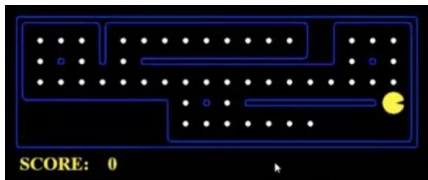
The *search state* keeps only the details needed for planning

## ■ Problem: Pathing

- States: (x, y) location
- Actions: NSEW
- Successor:

# What's in a State Space?

The *world state* includes every last detail of the environment



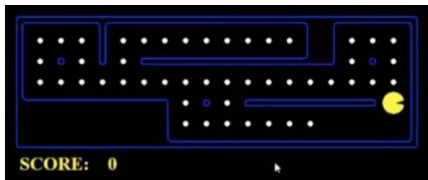
The *search state* keeps only the details needed for planning

## ■ Problem: Pathing

- States: (x, y) location
- Actions: NSEW
- Successor: update location only
- Goal test:

# What's in a State Space?

The *world state* includes every last detail of the environment



The *search state* keeps only the details needed for planning

## ■ Problem: Pathing

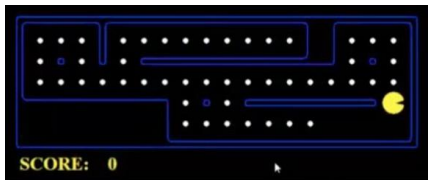
- States:  $(x, y)$  location
- Actions: NSEW
- Successor: update location only
- Goal test: is  $(x, y) = \text{END}$ ?

## ■ Problem: Eat-All-Dots

- States:

# What's in a State Space?

The *world state* includes every last detail of the environment



The *search state* keeps only the details needed for planning

## ■ Problem: Pathing

- States:  $(x, y)$  location
- Actions: NSEW
- Successor: update location only
- Goal test: is  $(x, y) = \text{END}$ ?

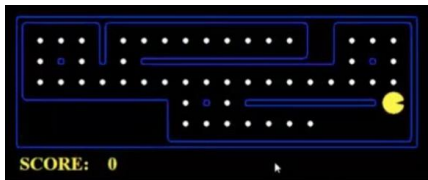
## ■ Problem: Eat-All-Dots

- States:  $\{(x, y), \text{dot booleans}\}$
- Actions:



# What's in a State Space?

The *world state* includes every last detail of the environment



The *search state* keeps only the details needed for planning

## ■ Problem: Pathing


- States:  $(x, y)$  location
- Actions: NSEW
- Successor: update location only
- Goal test: is  $(x, y) = \text{END}$ ?

## ■ Problem: Eat-All-Dots

- States:  $\{(x, y), \text{dot booleans}\}$
- Actions: NSEW
- Successor:

# What's in a State Space?

The *world state* includes every last detail of the environment



The *search state* keeps only the details needed for planning

## ■ Problem: Pathing

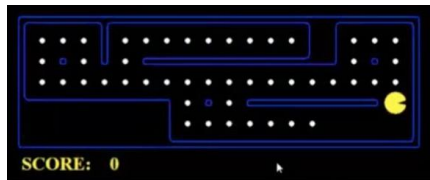
- States:  $(x, y)$  location
- Actions: NSEW
- Successor: update location only
- Goal test: is  $(x, y) = \text{END}$ ?

## ■ Problem: Eat-All-Dots

- States:  $\{(x, y), \text{dot booleans}\}$
- Actions: NSEW
- Successor: update location and possibly a dot boolean
- Goal test:

# What's in a State Space?

The *world state* includes every last detail of the environment



The *search state* keeps only the details needed for planning

## ■ Problem: Pathing

- States:  $(x, y)$  location
- Actions: NSEW
- Successor: update location only
- Goal test: is  $(x, y) = \text{END}$ ?

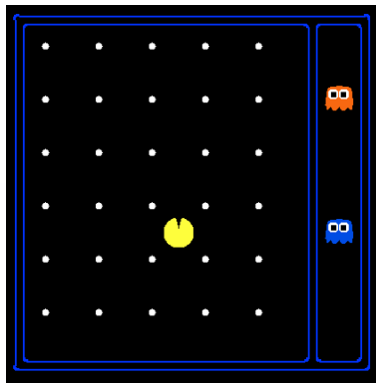
## ■ Problem: Eat-All-Dots

- States:  $\{(x, y), \text{dot booleans}\}$
- Actions: NSEW
- Successor: update location and possibly a dot boolean
- Goal test: dots all false

# State Space Sizes

## ■ Environment

- Agent positions: 120
- Food count: 30
- Ghost positions: 12
- Agent facing: NSEW



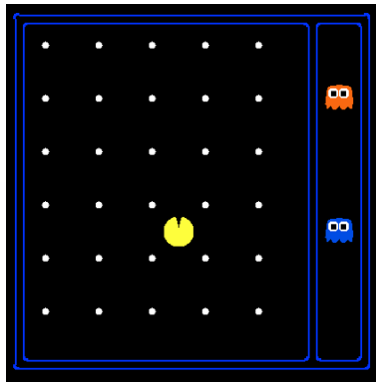
# State Space Sizes

## ■ Environment

- Agent positions: 120
- Food count: 30
- Ghost positions: 12
- Agent facing: NSEW

## ■ How many?

- World states?



# State Space Sizes

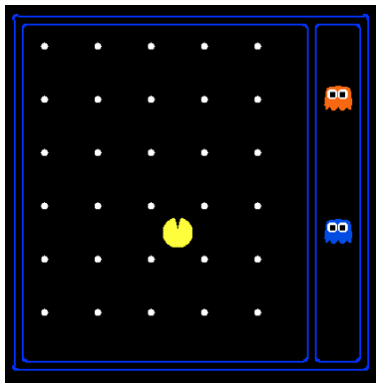
## ■ Environment

- Agent positions: 120
- Food count: 30
- Ghost positions: 12
- Agent facing: NSEW

## ■ How many?

- World states?

$$120 \times 2^{30} \times 12^2 \times 4$$



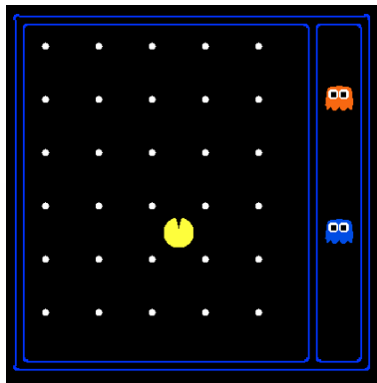
# State Space Sizes

## ■ Environment

- Agent positions: 120
- Food count: 30
- Ghost positions: 12
- Agent facing: NSEW

## ■ How many?

- World states?  
 $120 \times 2^{30} \times 12^2 \times 4$
- Search states? → Eat-All-Dots



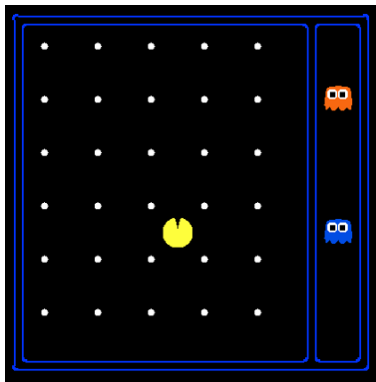
# State Space Sizes

## ■ Environment

- Agent positions: 120
- Food count: 30
- Ghost positions: 12
- Agent facing: NSEW

## ■ How many?

- World states?  
 $120 \times 2^{30} \times 12^2 \times 4$
- Search states? → Eat-All-Dots  
 $120 \times 2^{30}$





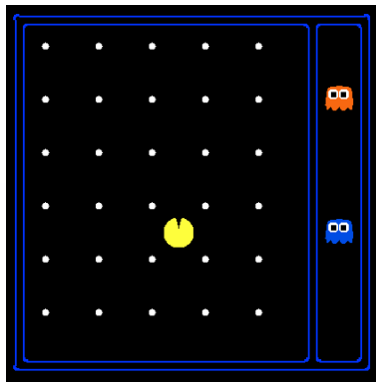
# State Space Sizes

## ■ Environment

- Agent positions: 120
- Food count: 30
- Ghost positions: 12
- Agent facing: NSEW

## ■ How many?

- World states?  
 $120 \times 2^{30} \times 12^2 \times 4$
- Search states? → Eat-All-Dots  
 $120 \times 2^{30}$
- Search states? → Pathing



# State Space Sizes

## ■ Environment

- Agent positions: 120
- Food count: 30
- Ghost positions: 12
- Agent facing: NSEW

## ■ How many?

### • World states?


$$120 \times 2^{30} \times 12^2 \times 4$$

### • Search states? → Eat-All-Dots

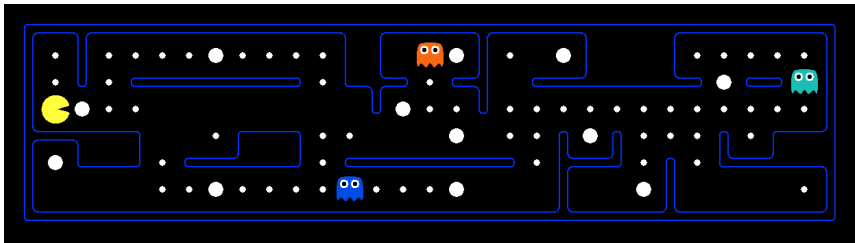
$$120 \times 2^{30}$$

### • Search states? → Pathing

$$120$$

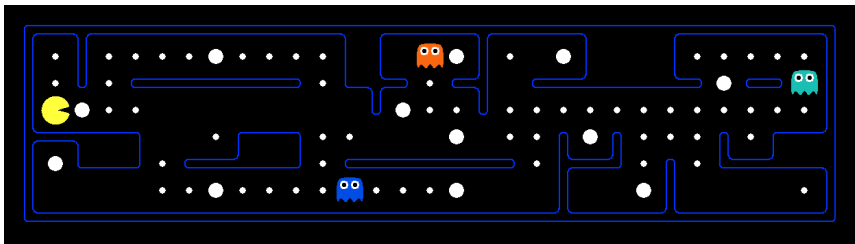


## Quiz: Safe Passage




- Problem: Eat all dots while keeping the ghosts perma-scared

## Quiz: Safe Passage



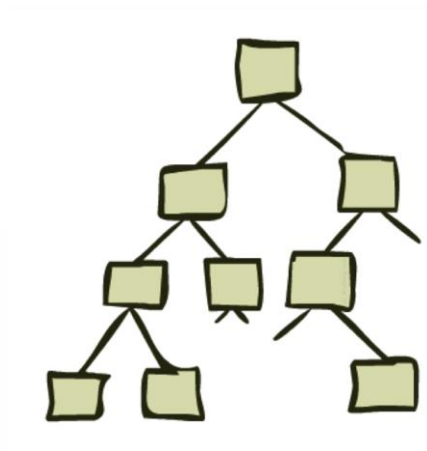
- Problem: Eat all dots while keeping the ghosts perma-scared
- State space?

## Quiz: Safe Passage




- Problem: Eat all dots while keeping the ghosts perma-scared
- State space?→ Agent position, dot booleans, power pellet booleans, remaining scared time

# State Space Graphs and Search Trees



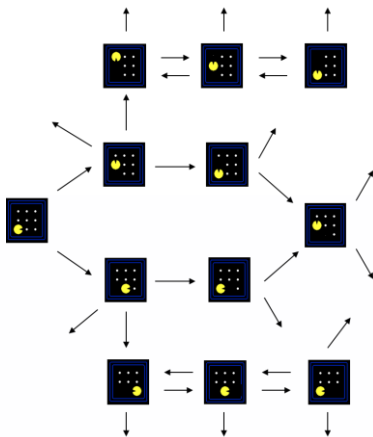
# State Space Graphs

- A mathematical representation of a search problem
  - Nodes are (abstracted) world configuration
  - Arcs represent successors (action results)
  - The goal test is a set of goal nodes (maybe only one)



# State Space Graphs

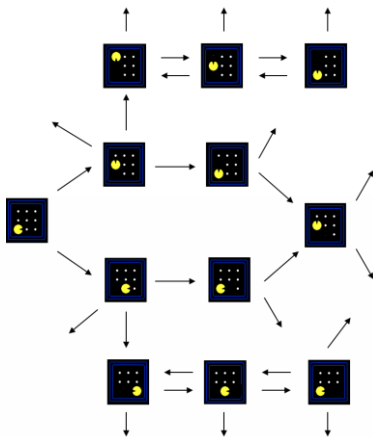
- A mathematical representation of a search problem
  - Nodes are (abstracted) world configuration
  - Arcs represent successors (action results)
  - The goal test is a set of goal nodes (maybe only one)
- Each state occurs only once






# State Space Graphs

- A mathematical representation of a search problem
  - Nodes are (abstracted) world configuration
  - Arcs represent successors (action results)
  - The goal test is a set of goal nodes (maybe only one)
- Each state occurs only once
- We can rarely build this full graph in memory

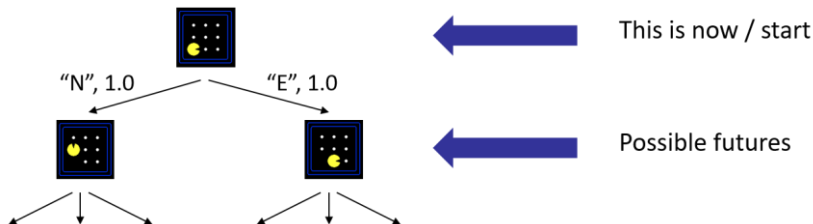


# State Space Graphs

- A mathematical representation of a search problem
  - Nodes are (abstracted) world configuration
  - Arcs represent successors (action results)
  - The goal test is a set of goal nodes (maybe only one)
- Each state occurs only once
- We can rarely build this full graph in memory




# Search Trees



- A "what if" tree of plans and their outcomes
- The start state is the root node
- Children correspond to successors
- Nodes show states, but correspond to PLANS that achieve those states
- For most problems, we can never actually build the tree

# State Space Graphs vs. Search Trees


## State Space Graph



Each **NODE** in in the search tree is an entire **PATH** in the state space graph.

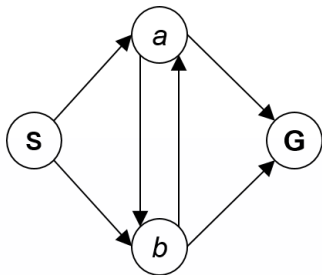
*We construct both  
on demand – and  
we construct as  
little as possible.*

## Search Tree



## Quiz: State Space Graphs vs. Search Trees

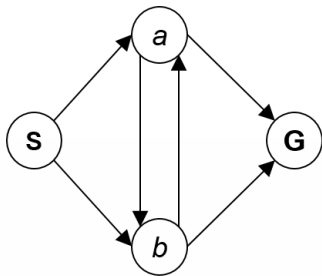
Consider this 4-state graph:



Let,  $s$  be the root.

## Quiz: State Space Graphs vs. Search Trees

Consider this 4-state graph:




Let,  $s$  be the root.

How big is the search tree?

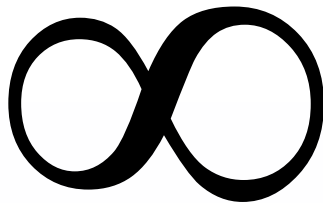
## Quiz: State Space Graphs vs. Search Trees

Consider this 4-state graph:



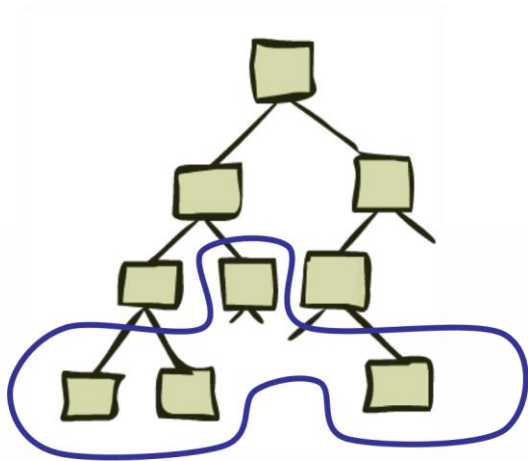
Let,  $s$  be the root.

How big is the search tree?




Lots of repeated structures in the search tree!

# Tree Search

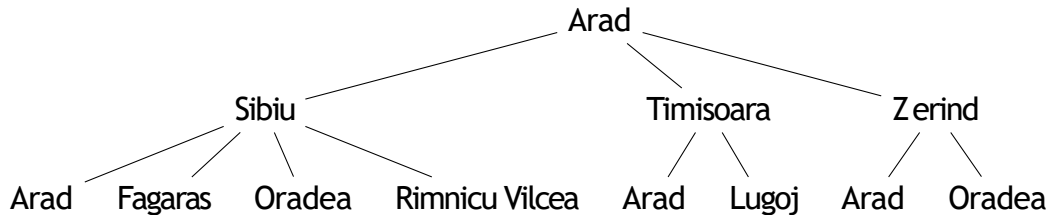




# Search Example: Romania



## Searching with a Search Tree



- Expand out potential plans
- Maintain a fringe (list) of partial plans under consideration
- Try to expand as few tree nodes as possible

# General Tree Search

function **TREE-SEARCH**(*problem*, *strategy*) returns a solution, or failure

# General Tree Search

function **TREE-SEARCH**(*problem*, *strategy*) returns a solution, or failure  
    initialize the search tree using the initial state of *problem*

# General Tree Search

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
```

# General Tree Search

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion
      then return failure
```

# General Tree Search

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion
      then return failure
    choose a leaf node according to strategy
```

# General Tree Search

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion
      then return failure
    choose a leaf node according to strategy
    if the node contains a goal state
      then return the corresponding solution
```



# General Tree Search

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion
      then return failure
    choose a leaf node according to strategy
    if the node contains a goal state
      then return the corresponding solution
    else
      expand the node and add the resulting node to the search tree
```

# General Tree Search

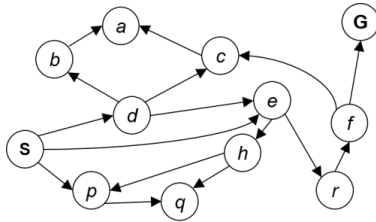
```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion
      then return failure
    choose a leaf node according to strategy
    if the node contains a goal state
      then return the corresponding solution
    else
      expand the node and add the resulting node to the search tree
  end
```

# General Tree Search

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion
      then return failure
    choose a leaf node according to strategy
    if the node contains a goal state
      then return the corresponding solution
    else
      expand the node and add the resulting node to the search tree
  end
```

- Concepts: Fringe, Expansion, Exploration strategy

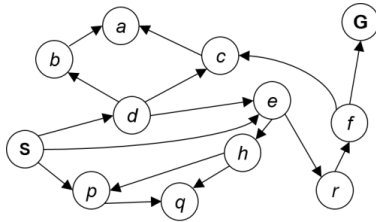
# Example: Tree Search



Search Tree:

Fringe:

# Example: Tree Search



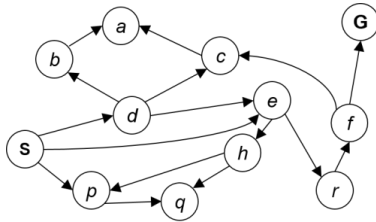
Search Tree:



Fringe:

s

# Example: Tree Search



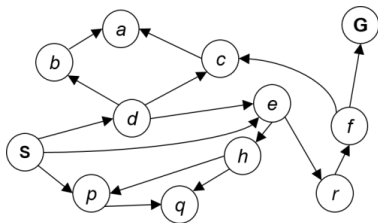
Search Tree:



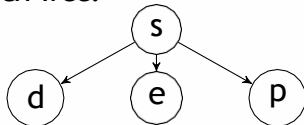
Fringe:

s

# Example: Tree Search




Search Tree:




Fringe:

$s \rightarrow d, s \rightarrow e, s \rightarrow p$

# Example: Tree Search



Search Tree:

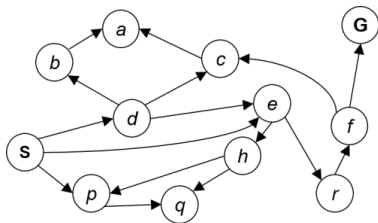


Fringe:

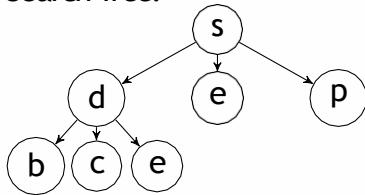
~~s~~ → d, s → e, s → p



# Example: Tree Search



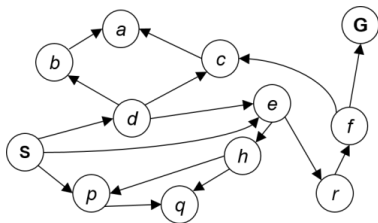
Search Tree:



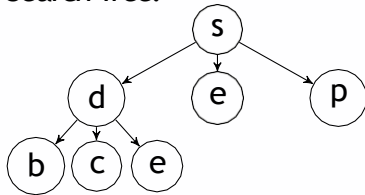
Fringe:

$s \rightarrow e, s \rightarrow p,$   
 $s \rightarrow d \rightarrow b, s \rightarrow d \rightarrow c, s \rightarrow d \rightarrow$   
 $e$

# Example: Tree Search




Search Tree:




Fringe:

$s \rightarrow e, s \rightarrow p,$   
 $s \rightarrow d \rightarrow b, s \rightarrow d \rightarrow c, s \rightarrow d \rightarrow$   
 $e$

# Example: Tree Search



Search Tree:



Fringe:


$s \rightarrow e, s \rightarrow p,$   
 $s \rightarrow d \rightarrow c, s \rightarrow d \rightarrow e,$   
 $s \rightarrow d \rightarrow b \rightarrow a \dots$

# Depth-First Search

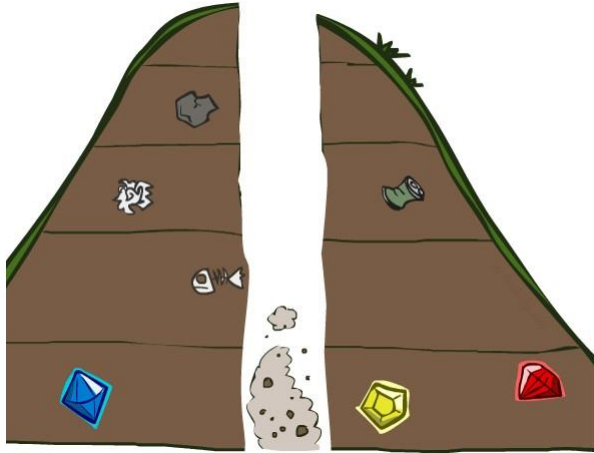


# Depth-First Search

- Strategy: Expand the deepest node first
- Implementation: Fringe is a LIFO stack



# Search Algorithm Properties



# Search Algorithm Properties

- Complete: Guaranteed to find solution if one exists?
- Optimal: Guaranteed to find the least cost path?

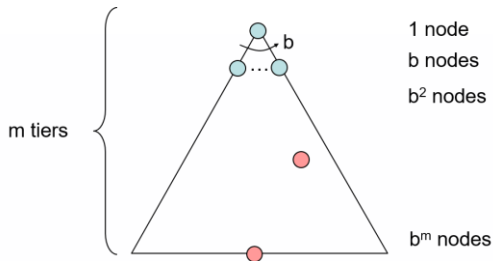
# Search Algorithm Properties

- Complete: Guaranteed to find solution if one exists?
- Optimal: Guaranteed to find the least cost path?
- Time complexity?
- Space complexity?



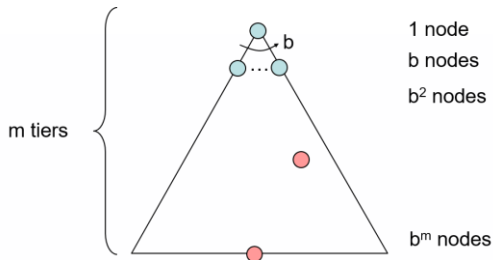
# Search Algorithm Properties

- Complete: Guaranteed to find solution if one exists?
- Optimal: Guaranteed to find the least cost path?
- Time complexity?
- Space complexity?
- Cartoon of a search tree:
  - $b$  is the branching factor
  - $m$  is the maximum depth
  - solutions at various depths



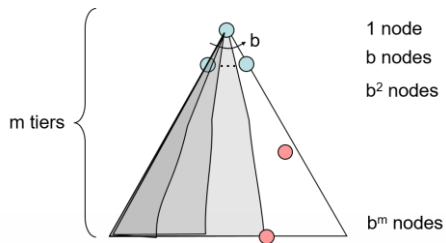
# Search Algorithm Properties

- Complete: Guaranteed to find solution if one exists?
- Optimal: Guaranteed to find the least cost path?
- Time complexity?
- Space complexity?
- Cartoon of a search tree:
  - $b$  is the branching factor
  - $m$  is the maximum depth
  - solutions at various depths
- Number of nodes in entire tree?
  - $1 + b + b^2 + \dots + b^m = O(b^m)$



# Depth-First Search (DFS) Properties

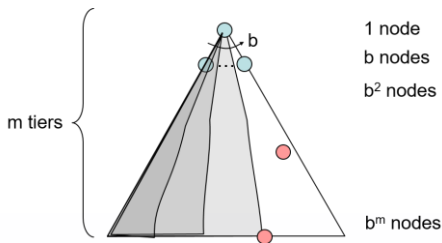
- What nodes does DFS expand?



# Depth-First Search (DFS) Properties

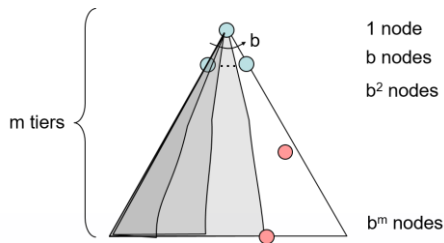
## ■ What nodes does DFS expand?

- Some left prefix of the tree
- Could process the whole tree
- If  $m$  is finite, takes time  $O(b^m)$



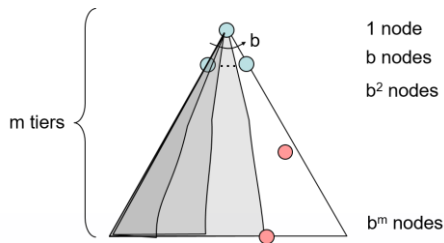
# Depth-First Search (DFS) Properties

- What nodes does DFS expand?
  - Some left prefix of the tree
  - Could process the whole tree
  - If  $m$  is finite, takes time  $O(b^m)$
- How much does a fringe take?



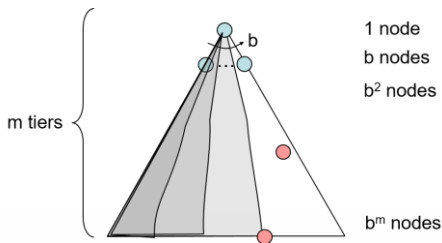
# Depth-First Search (DFS) Properties

- What nodes does DFS expand?
  - Some left prefix of the tree
  - Could process the whole tree
  - If  $m$  is finite, takes time  $O(b^m)$
- How much does a fringe take?
  - Only has siblings on path to root,  $O(bm)$



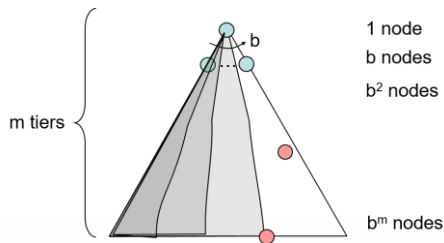
# Depth-First Search (DFS) Properties

- What nodes does DFS expand?
  - Some left prefix of the tree
  - Could process the whole tree
  - If  $m$  is finite, takes time  $O(b^m)$
- How much does a fringe take?
  - Only has siblings on path to root,  $O(bm)$
- Is it complete?



# Depth-First Search (DFS) Properties

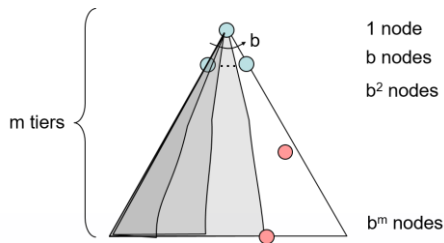
- What nodes does DFS expand?
  - Some left prefix of the tree
  - Could process the whole tree
  - If  $m$  is finite, takes time  $O(b^m)$
- How much does a fringe take?
  - Only has siblings on path to root,  $O(bm)$
- Is it complete?
  - $m$  could be infinite, so only if we prevent cycles





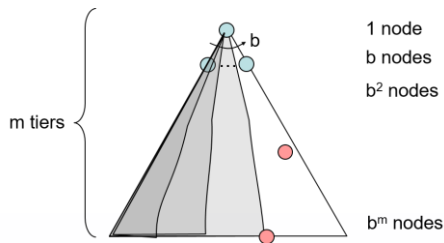
# Depth-First Search (DFS) Properties

- What nodes does DFS expand?
  - Some left prefix of the tree
  - Could process the whole tree
  - If  $m$  is finite, takes time  $O(b^m)$
- How much does a fringe take?
  - Only has siblings on path to root,  $O(bm)$
- Is it complete?
  - $m$  could be infinite, so only if we prevent cycles
- Is it optimal?




# Depth-First Search (DFS) Properties

- What nodes does DFS expand?
  - Some left prefix of the tree
  - Could process the whole tree
  - If  $m$  is finite, takes time  $O(b^m)$
- How much does a fringe take?
  - Only has siblings on path to root,  $O(bm)$
- Is it complete?
  - $m$  could be infinite, so only if we prevent cycles
- Is it optimal?
  - No, it finds the "leftmost" solution, regardless of the depth/cost

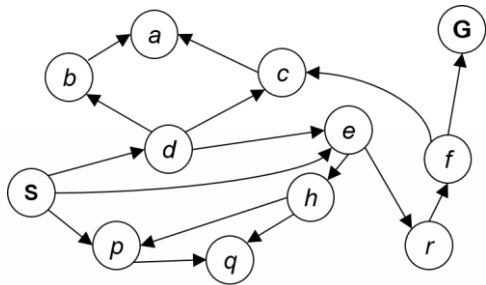


# Breadth-First Search



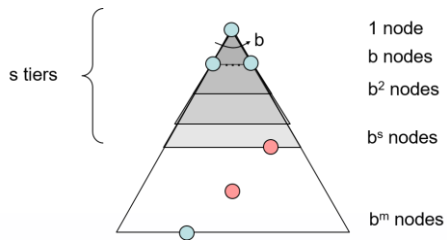
# Breadth-First Search

- Strategy: Expand the shallowest node first
- Implementation: Fringe is a FIFO queue



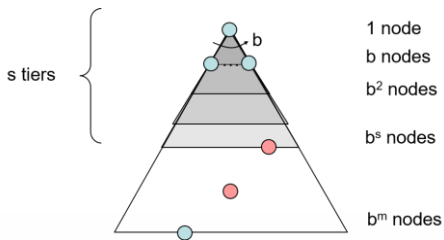
# Breadth-First Search (BFS) Properties

- What nodes does BFS expand?



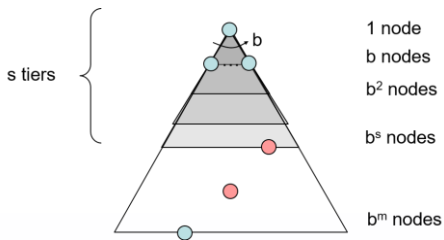
# Breadth-First Search (BFS) Properties

- What nodes does BFS expand?
  - Process all nodes above shallowest solution
  - Let depth of the shallowest solution be  $s$
  - Search takes time  $O(b^s)$



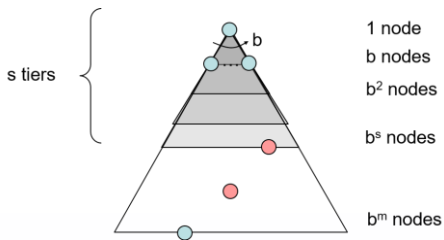
# Breadth-First Search (BFS) Properties

- What nodes does BFS expand?
  - Process all nodes above shallowest solution
  - Let depth of the shallowest solution be  $s$
  - Search takes time  $O(b^s)$
- How much does a fringe take?



# Breadth-First Search (BFS) Properties

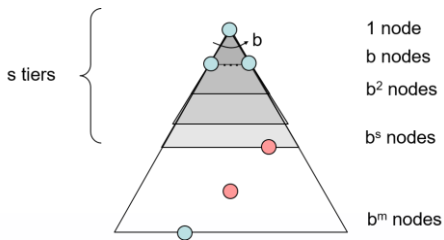
- What nodes does BFS expand?
  - Process all nodes above shallowest solution
  - Let depth of the shallowest solution be  $s$
  - Search takes time  $O(b^s)$
- How much does a fringe take?
  - Has roughly the last tier,  $O(b^s)$






# Breadth-First Search (BFS) Properties

- What nodes does BFS expand?
  - Process all nodes above shallowest solution
  - Let depth of the shallowest solution be  $s$
  - Search takes time  $O(b^s)$
- How much does a fringe take?
  - Has roughly the last tier,  $O(b^s)$
- Is it complete?



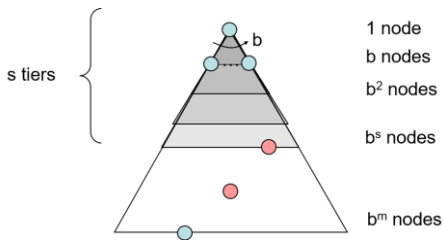
# Breadth-First Search (BFS) Properties

- What nodes does BFS expand?
  - Process all nodes above shallowest solution
  - Let depth of the shallowest solution be  $s$
  - Search takes time  $O(b^s)$
- How much does a fringe take?
  - Has roughly the last tier,  $O(b^s)$
- Is it complete?
  - $s$  must be finite if a solution exists, yes!



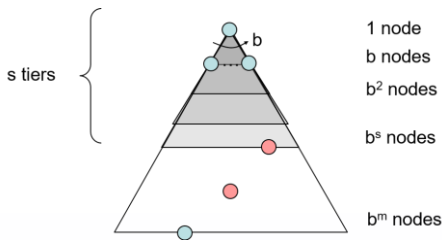
# Breadth-First Search (BFS) Properties

- What nodes does BFS expand?
  - Process all nodes above shallowest solution
  - Let depth of the shallowest solution be  $s$
  - Search takes time  $O(b^s)$
- How much does a fringe take?
  - Has roughly the last tier,  $O(b^s)$
- Is it complete?
  - $s$  must be finite if a solution exists, yes!
- Is it optimal?



# Breadth-First Search (BFS) Properties

- What nodes does BFS expand?
  - Process all nodes above shallowest solution
  - Let depth of the shallowest solution be  $s$
  - Search takes time  $O(b^s)$
- How much does a fringe take?
  - Has roughly the last tier,  $O(b^s)$
- Is it complete?
  - $s$  must be finite if a solution exists, yes!
- Is it optimal?
  - Only if costs are all 1



# Quiz: DFS vs BFS



Video: [Empty-BFS](#), [Empty-DFS](#), [MazeWater-BFS](#), [MazeWater-DFS](#)

# Quiz: DFS vs BFS

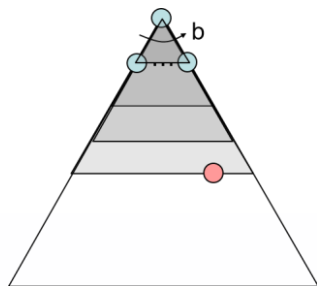


- When will BFS outperform DFS?
- When will DFS outperform BFS?

Video: [Empty-BFS](#), [Empty-DFS](#), [MazeWater-BFS](#), [MazeWater-DFS](#)


# Iterative Deepening

- Idea: get DFS's space advantage with BFS's time / shallow-solution advantages



# Iterative Deepening


- Idea: get DFS's space advantage with BFS's time / shallow-solution advantages
  - Run DFS with depth limit 1. If no solution...
  - Run DFS with depth limit 2. If no solution...
  - Run DFS with depth limit 3. ...



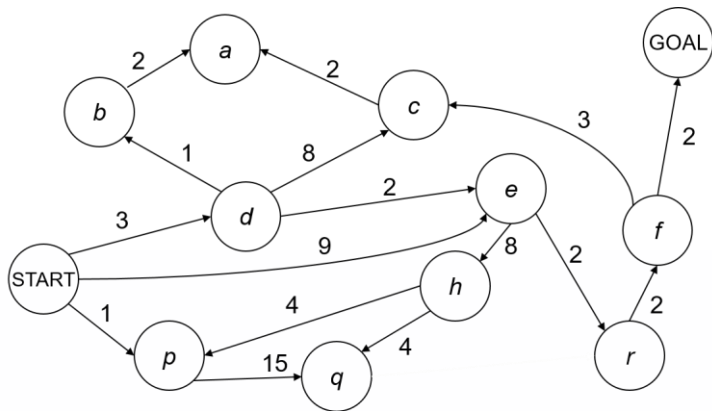


# Iterative Deepening

- Idea: get DFS's space advantage with BFS's time / shallow-solution advantages
  - Run DFS with depth limit 1. If no solution...
  - Run DFS with depth limit 2. If no solution...
  - Run DFS with depth limit 3. ...
- Isn't that wastefully redundant?
  - Generally most work happens in the lowest level searched, so not so bad!

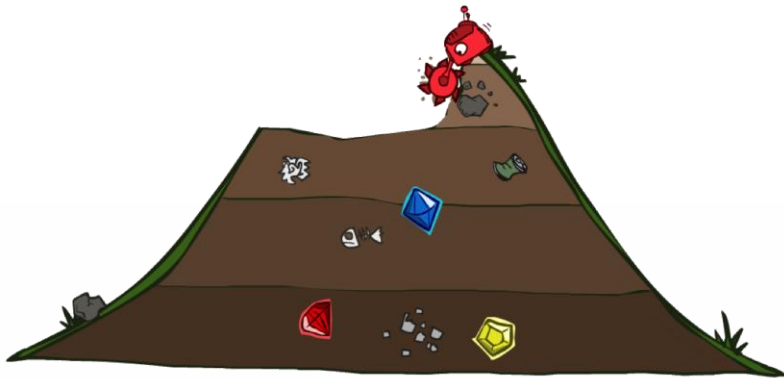


# Cost-Sensitive Search



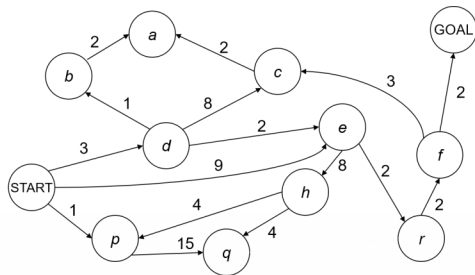
BFS is optimal in terms of the number of actions performed.  
It does not find the least cost path.

# Uniform Cost Search



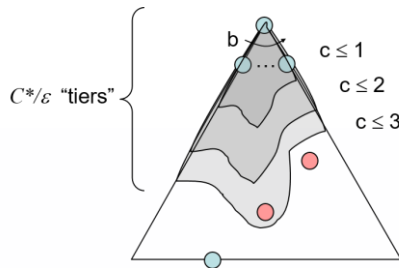
# Uniform Cost Search

- Strategy: Expand the cheapest node first
- Implementation: Fringe is a priority queue
  - Priority: Cumulative cost




# Uniform Cost Search (UCS) Properties

- What nodes does UCS expand?



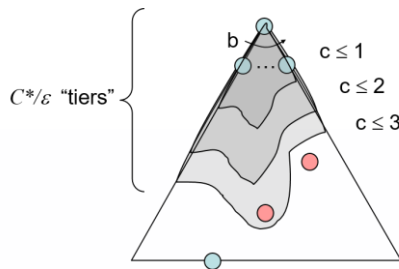
# Uniform Cost Search (UCS) Properties

- What nodes does UCS expand?
  - Process all nodes with cost less than cheapest solution
  - If the solution costs  $C^*$  and arcs cost at least  $\epsilon$ , then "effective depth" is roughly  $C^*/\epsilon$
  - Takes time  $O(b^{C^*/\epsilon})$



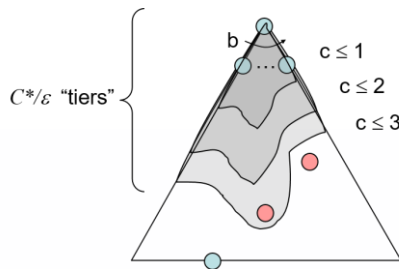
# Uniform Cost Search (UCS) Properties

- What nodes does UCS expand?
  - Process all nodes with cost less than cheapest solution
  - If the solution costs  $C^*$  and arcs cost at least  $\epsilon$ , then "effective depth" is roughly  $C^*/\epsilon$
  - Takes time  $O(b^{C^*/\epsilon})$
- How much does a fringe take?



# Uniform Cost Search (UCS) Properties


- What nodes does UCS expand?
  - Process all nodes with cost less than cheapest solution
  - If the solution costs  $C^*$  and arcs cost at least  $\epsilon$ , then "effective depth" is roughly  $C^*/\epsilon$
  - Takes time  $O(b^{C^*/\epsilon})$
- How much does a fringe take?
  - Has roughly the last tier,  $O(b^{C^*/\epsilon})$





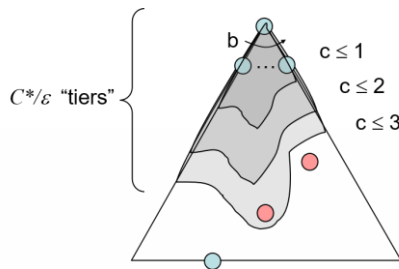
# Uniform Cost Search (UCS) Properties

- What nodes does UCS expand?
  - Process all nodes with cost less than cheapest solution
  - If the solution costs  $C^*$  and arcs cost at least  $\epsilon$ , then "effective depth" is roughly  $C^*/\epsilon$
  - Takes time  $O(b^{C^*/\epsilon})$
- How much does a fringe take?
  - Has roughly the last tier,  $O(b^{C^*/\epsilon})$
- Is it complete?




# Uniform Cost Search (UCS) Properties

- What nodes does UCS expand?
  - Process all nodes with cost less than cheapest solution
  - If the solution costs  $C^*$  and arcs cost at least  $\epsilon$ , then "effective depth" is roughly  $C^*/\epsilon$
  - Takes time  $O(b^{C^*/\epsilon})$
- How much does a fringe take?
  - Has roughly the last tier,  $O(b^{C^*/\epsilon})$
- Is it complete?
  - Assuming best solution has a finite cost and minimum arc cost is positive, yes



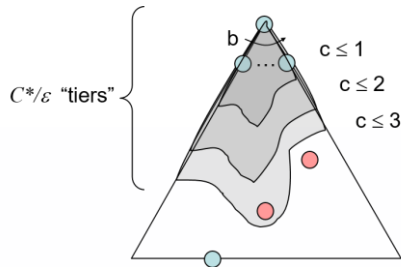
# Uniform Cost Search (UCS) Properties

- What nodes does UCS expand?
  - Process all nodes with cost less than cheapest solution
  - If the solution costs  $C^*$  and arcs cost at least  $\epsilon$ , then "effective depth" is roughly  $C^*/\epsilon$
  - Takes time  $O(b^{C^*/\epsilon})$
- How much does a fringe take?
  - Has roughly the last tier,  $O(b^{C^*/\epsilon})$
- Is it complete?
  - Assuming best solution has a finite cost and minimum arc cost is positive, yes
- Is it optimal?



# Uniform Cost Search (UCS) Properties

- What nodes does UCS expand?
  - Process all nodes with cost less than cheapest solution
  - If the solution costs  $C^*$  and arcs cost at least  $\epsilon$ , then "effective depth" is roughly  $C^*/\epsilon$
  - Takes time  $O(b^{C^*/\epsilon})$
- How much does a fringe take?
  - Has roughly the last tier,  $O(b^{C^*/\epsilon})$
- Is it complete?
  - Assuming best solution has a finite cost and minimum arc cost is positive, yes
- Is it optimal?
  - Yes! (Proof, in future class)




# Uniform Cost Issues

- Remember: UCS explores increasing cost contours
- The good: UCS is complete and optimal

Video: [MazeShallowDeep-UCS](#), [MazeShallowDeep-BFS](#), [MazeShallowDeep-DFS](#)

# Uniform Cost Issues

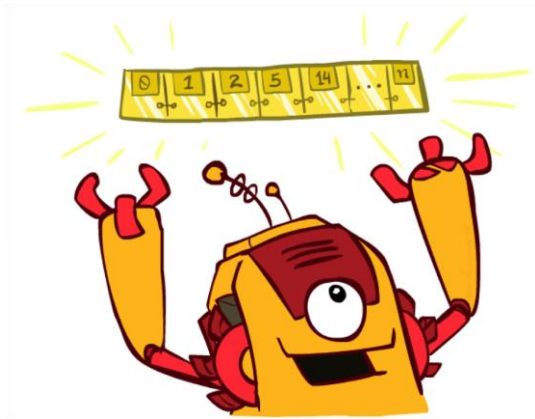
- Remember: UCS explores increasing cost contours
- The good: UCS is complete and optimal
- The bad
  - Explores options in every "direction"
  - No information about goal location



Video: [MazeShallowDeep-UCS](#), [MazeShallowDeep-BFS](#), [MazeShallowDeep-DFS](#)


# The One Queue

- All these search algorithms are the same except for fringe strategies
  - Conceptually, all fringes are priority queues (i.e. collections of nodes with attached priorities)
  - Practically, for DFS and BFS, you can avoid the  $\log(n)$  overhead from an actual priority queue, by using stacks and queues
  - Can even code one implementation that takes a variable queuing object



# Search and Models

- Search operates over models of the world
  - The agent doesn't actually try all the plans out in the real world!
  - Planning is all "in simulation"
  - Your search is only as good as your models...





## Suggested Reading

- Russell & Norvig: Chapter: 3.1-3.4
- Poole & Mackworth: Chapter: 3.1-3.5