


# Artificial Intelligence

## CSE 4617

Ahnaf Munir  
Assistant Professor  
Islamic University of Technology

# Uncertain Outcomes



# Recap: Probabilities

- **Random variable** → Event whose outcome is unknown
- **Probability distribution** → Assignment of weights to outcomes
- **Example: Traffic on freeway**
  - Random variable:  $T$  = whether there's traffic
  - Outcome:  $T \in \{\text{none, light, heavy}\}$
  - Distribution:  $P(T = \text{none}) = 0.25$ ,  $P(T = \text{light}) = 0.50$ ,  $P(T = \text{heavy}) = 0.25$
- **Some laws of probability (more later):**
  - Non-negative
  - Sum of probabilities over all possible outcomes: 1
- **As we get more evidence, probabilities may change:**
  - $P(T = \text{heavy}) = 0.25$ ,  $P(T = \text{heavy} | H = 8 \text{ a.m.}) = 0.60$



0.25



0.50



0.25

# Recap: Expectations

- Expected value of a function of random variable
- Average, weighted by the probability distribution over outcomes



# Recap: Expectations


- Expected value of a function of random variable
- Average, weighted by the probability distribution over outcomes
- Example: How long to get to the airport?





Time:	20 min	30 min	60 min
Probability:	0.25	0.50	0.25





# Worst-Case vs. Average Case



# Worst-Case vs. Average Case





# Worst-Case vs. Average Case







# Worst-Case vs. Average Case



Idea: Uncertain outcomes controlled by chance, not an adversary!


# Worst-Case vs. Average Case



Idea: Uncertain outcomes controlled by chance, not an adversary!

# Expectimax Search


- Why wouldn't we know what the result of an action will be?



Video: [minimax](#), [expectimax](#)

# Expectimax Search


- Why wouldn't we know what the result of an action will be?
  - Explicit randomness: rolling dice



Video: [minimax](#), [expectimax](#)

# Expectimax Search


- Why wouldn't we know what the result of an action will be?
  - Explicit randomness: rolling dice
  - Unpredictable opponents: the ghosts respond randomly



Video: [minimax](#), [expectimax](#)

# Expectimax Search


- Why wouldn't we know what the result of an action will be?
  - Explicit randomness: rolling dice
  - Unpredictable opponents: the ghosts respond randomly
  - Failed actions: when moving a robot, wheels might slip



Video: [minimax](#), [expectimax](#)

# Expectimax Search


- Why wouldn't we know what the result of an action will be?
  - Explicit randomness: rolling dice
  - Unpredictable opponents: the ghosts respond randomly
  - Failed actions: when moving a robot, wheels might slip
- Values should now reflect average-case (expectimax) outcomes, not worst-case (minimax) outcome



Video: [minimax](#), [expectimax](#)

# Expectimax Search

- Why wouldn't we know what the result of an action will be?
  - Explicit randomness: rolling dice
  - Unpredictable opponents: the ghosts respond randomly
  - Failed actions: when moving a robot, wheels might slip
- Values should now reflect average-case (expectimax) outcomes, not worst-case (minimax) outcome
- Compute the average score under optimal play
  - Max nodes as in minimax search
  - Chance nodes are like min nodes but the outcome is uncertain
  - Calculate their **expected utilities**  
i.e. take weighted average (expectation) of children



Video: [minimax](#), [expectimax](#)



# Expectimax Pseudocode

def *value*(*state*):

    if the state is a terminal state: return the state's utility

    if the next agent is MAX: return *max-value*(*state*)

    if the next agent is EXP: return *exp-value*(*state*)

def *max-value*(*state*):

    initialize  $v = -\infty$

    for each successor of *state*:

$v = \max(v, \textit{value}(\textit{successor}))$

    return  $v$

def *exp-value*(*state*):

    initialize  $v = 0$

    for each successor of *state*:

$p = \textit{probability}(\textit{successor})$

$v += p \times \textit{value}(\textit{successor})$

    return  $v$

# Expectimax Pseudocode

def exp-value(*state*):


    initialize  $v = 0$

    for each successor of *state*:


$p = \text{probability}(\text{successor})$

$v += p \times \text{value}(\text{successor})$


    return  $v$





# Expectimax Quiz




# Expectimax Pruning?




# Expectimax Pruning?




## Depth-Limited Expectimax



# Depth-Limited Expectimax




# Depth-Limited Expectimax






# Probabilities



# What Probabilities to Use?

- In expectimax search, we have a probabilistic model of how the opponent (or environment) will behave in any state
  - Model could be a simple uniform distribution (roll a die)
  - Model could be sophisticated and require a great deal of computation
  - We have a chance node for any outcome out of our control: opponent or environment
  - The model might say that adversarial actions are likely!
- For now, assume each chance node magically comes along with probabilities that specify the distribution over its outcomes



*Having a probabilistic belief about another agent's action does not mean that the agent is flipping any coins!*

# Informed Probabilities

- Let's say you know that your opponent is actually running a depth 2 minimax, using the result 80% of the time, and moving randomly otherwise
- Question: What tree search should you use?


# Informed Probabilities

- Let's say you know that your opponent is actually running a depth 2 minimax, using the result 80% of the time, and moving randomly otherwise
- Question: What tree search should you use?
  - Answer: Expectimax!

# Informed Probabilities

- Let's say you know that your opponent is actually running a depth 2 minimax, using the result 80% of the time, and moving randomly otherwise
- Question: What tree search should you use?


- Answer: Expectimax!
  - To figure out EACH chance node's probabilities, you have to run a simulation of your opponent



# Informed Probabilities


- Let's say you know that your opponent is actually running a depth 2 minimax, using the result 80% of the time, and moving randomly otherwise
- Question: What tree search should you use?

- Answer: Expectimax!
  - To figure out EACH chance node's probabilities, you have to run a simulation of your opponent



# Informed Probabilities


- Let's say you know that your opponent is actually running a depth 2 minimax, using the result 80% of the time, and moving randomly otherwise
- Question: What tree search should you use?



- Answer: Expectimax!
  - To figure out EACH chance node's probabilities, you have to run a simulation of your opponent
  - This kind of things gets very slow very quickly
  - Even worse if you have to simulate your opponent simulating you...

# Informed Probabilities


- Let's say you know that your opponent is actually running a depth 2 minimax, using the result 80% of the time, and moving randomly otherwise
- Question: What tree search should you use?



- Answer: Expectimax!
  - To figure out EACH chance node's probabilities, you have to run a simulation of your opponent
  - This kind of things gets very slow very quickly
  - Even worse if you have to simulate your opponent simulating you...
  - ...except for minimax, which has the nice property that it all collapses into one game tree.



# Modeling Assumptions



# The Dangers of Optimism and Pessimism

# The Dangers of Optimism and Pessimism

## Dangerous Optimism

Assuming chance when the world is  
adversarial



# The Dangers of Optimism and Pessimism

## Dangerous Optimism

Assuming chance when the world is adversarial




## Dangerous Pessimism

Assuming the worst case when it's not likely



# Assumptions vs. Reality




	Adversarial Ghost	Random Ghost
Minimax Pacman		
Expectimax Pacman		

Pacman used depth 4 search with an eval function that avoids trouble  
Ghost used depth 2 search with an eval function that seeks Pacman

Videos: [randGhostExpPac](#), [advGhostMiniPac](#), [miniGhostExpPac](#), [randGhostMiniPac](#)

# Assumptions vs. Reality



	Adversarial Ghost	Random Ghost
Minimax Pacman	Won 5/5 Avg. Score: 483	Won 5/5 Avg. Score: 493
Expectimax Pacman	Won 1/5 Avg. Score: -303	Won 5/5 Avg. Score: 503


Results from playing 5 games

Pacman used depth 4 search with an eval function that avoids trouble

Ghost used depth 2 search with an eval function that seeks Pacman


Videos: [randGhostExpPac](#), [advGhostMiniPac](#), [miniGhostExpPac](#), [randGhostMiniPac](#)

# Other Game Types



# Mixed Layer Types

- E.g. Backgammon
- Expectiminimax
  - Environment is an extra “random agent” player that moves after each min/max agent
  - Each node computes the appropriate combination of its children






# Example: Backgammon

- Dice rolls increase  $b$ : 21 possible rolls with 2 dice
  - Backgammon  $\approx$  20 legal moves
  - Depth 2 =  $20 \times (21 \times 20)^3 = 1.2 \times 10^9$
- As depth increases, probability of reaching a given search node shrinks
  - Usefulness of search is diminished
  - Limiting depth is less damaging
  - Pruning is trickier...
- Historic AI: TDGammon uses depth-2 search + very good evaluation function + reinforcement learning: world-champion level play
- 1<sup>st</sup> AI world champion in any game!




# Multi-Agent Utilities

- What if the game is not zero-sum, or has multiple players?
- Generalization of minimax:
  - Terminals have utility tuples
  - Node values are also utility tuples
  - Each player maximizes its own component
  - Can give rise to cooperation and competition dynamically...



# Utilities



# Maximum Expected Utility

- Why should we average utilities? Why not minimax?

# Maximum Expected Utility

- Why should we average utilities? Why not minimax?




# Maximum Expected Utility

- Why should we average utilities? Why not minimax?
- Principle of maximum expected utility:
  - A rational agent should chose the action that maximizes its expected utility, given its knowledge


# Maximum Expected Utility

- Why should we average utilities? Why not minimax?
- Principle of maximum expected utility:
  - A rational agent should choose the action that maximizes its expected utility, given its knowledge
- Questions:
  - Where do utilities come from?
  - How do we know such utilities even exist?
  - How do we know that averaging even makes sense?
  - What if our behavior (preferences) can't be described by utilities?

# What Utilities to Use?




$x^2$






# What Utilities to Use?




- For worst-case minimax reasoning, terminal function scale doesn't matter
  - We just want better states to have higher evaluations (get the ordering right)
  - We call this **insensitivity to monotonic transformations**

# What Utilities to Use?



- For worst-case minimax reasoning, terminal function scale doesn't matter
  - We just want better states to have higher evaluations (get the ordering right)
  - We call this **insensitivity to monotonic transformations**

# What Utilities to Use?



- For worst-case minimax reasoning, terminal function scale doesn't matter
  - We just want better states to have higher evaluations (get the ordering right)
  - We call this **insensitivity to monotonic transformations**
- For average-case expectimax reasoning, we need *magnitudes* to be meaningful

# Utilities (Revisited)


- Utilities are functions from outcomes (states of the world) to real numbers that describe an agent's preferences
- Where do utilities come from?
  - In a game, may be simple (+1/-1)
  - Utilities summarize the agent's goals
  - Theorem: any "rational" preferences can be summarized as a utility function




# Utilities: Uncertain Outcomes




# Utilities: Uncertain Outcomes



# Utilities: Uncertain Outcomes



# Utilities: Uncertain Outcomes






# Preferences

- An agent must have preferences among:
  - Prizes:  $A, B$ , etc.
  - Lotteries: Situations with uncertain prizes  
 $L = [p, A; (1 - p), B]$
- Notation:
  - Preference:  $A \succ B$
  - Indifference:  $A \sim B$


A Prize



A Lottery



# Rationality



# Rational Preferences

- We want some constraints on preferences before we call them rational, such as:


Axiom of Transitivity:  $(A \succ B) \wedge (B \succ C) \Rightarrow (A \succ C)$

# Rational Preferences

- We want some constraints on preferences before we call them rational, such as:

Axiom of Transitivity:  $(A \succ B) \wedge (B \succ C) \Rightarrow (A \succ C)$

- For example: an agent with **intransitive preferences** can be induced to give away all of its money
  - If  $B \succ C$ , then an agent with  $C$  would pay (say) 1 cent to get  $B$
  - If  $A \succ B$ , then an agent with  $B$  would pay (say) 1 cent to get  $A$
  - If  $C \succ A$ , then an agent with  $A$  would pay (say) 1 cent to get  $C$



# Rational Preferences

## The Axioms of Rationality

- Orderability

$$(A \succ B) \vee (B \succ A) \vee (A \sim B)$$

- Transitivity

$$(A \succ B) \wedge (B \succ C) \Rightarrow (A \succ C)$$

- Continuity

$$A \succ B \succ C \Rightarrow \exists p [p, A; 1 - p, C] \sim B$$

- Substitutability

$$A \sim B \Rightarrow [p, A; 1 - p, C] \sim [p, B; 1 - p, C]$$

- Monotonicity

$$A \succ B \Rightarrow (p \geq q \Leftrightarrow [p, A; 1 - p, B] \geq [q, A; 1 - q, B])$$

**Theorem:** Rational preferences imply behavior describable as maximization of expected utility



# MEU Principle

- Theorem [Ramsey, 1931; von Neumann & Morgenstern, 1944]
  - Given any preferences satisfying these constraints, there exists a real-valued function  $U$  such that:  
$$U(A) \geq U(B) \Leftrightarrow A \geq B$$
$$U([p_1, S_1; \dots; p_n, S_n]) = \sum_i p_i U(S_i)$$
  - i.e. values assigned by  $U$  preserve preferences of both prizes and lotteries!



# MEU Principle

- Theorem [Ramsey, 1931; von Neumann & Morgenstern, 1944]

- Given any preferences satisfying these constraints, there exists a real-valued function  $U$  such that:

$$U(A) \geq U(B) \Leftrightarrow A \succeq B$$

$$U([p_1, S_1; \dots; p_n, S_n]) = \sum_i p_i U(S_i)$$


- i.e. values assigned by  $U$  preserve preferences of both prizes and lotteries!

- Maximum Expected Utility (MEU) principle:

- Choose the action that maximizes expected utility
- Note: an agent can be entirely rational (consistent with MEU) without ever representing or manipulating utilities and probabilities
- E.g., a lookup table for perfect tic-tac-toe, a reflex vacuum cleaner



# Human Utilities





# Utility Scales

- **Normalized utilities:**  $u_+ = 1.0, u_- = 0.0$

# Utility Scales

- **Normalized utilities:**  $u_+ = 1.0, u_- = 0.0$
- **Micromorts:** one-millionth chance of death, useful for paying to reduce product risks, etc.



# Utility Scales

- **Normalized utilities:**  $u_+ = 1.0, u_- = 0.0$
- **Micromorts:** one-millionth chance of death, useful for paying to reduce product risks, etc.
- **QALYs:** Quality-Adjusted Life Years, useful for medical decisions involving substantial risk



# Utility Scales

- **Normalized utilities:**  $u_+ = 1.0, u_- = 0.0$
- **Micromorts:** one-millionth chance of death, useful for paying to reduce product risks, etc.
- **QALYs:** Quality-Adjusted Life Years, useful for medical decisions involving substantial risk
- **Note:** behavior is invariant under positive linear transformation

$$U'(x) = k_1 U(x) + k_2 \text{ where } k_1 > 0$$




# Utility Scales

- **Normalized utilities:**  $u_+ = 1.0, u_- = 0.0$
- **Micromorts:** one-millionth chance of death, useful for paying to reduce product risks, etc.
- **QALYs:** Quality-Adjusted Life Years, useful for medical decisions involving substantial risk
- Note: behavior is invariant under positive linear transformation  
 $U'(x) = k_1 U(x) + k_2$  where  $k_1 > 0$
- With deterministic prizes only (no lottery choices), only **ordinal utility** can be determined, i.e., total order on prizes




# Human Utilities (Revisited)

- Utilities map states to real numbers



# Human Utilities (Revisited)

- Utilities map states to real numbers
- Standard approach to assessment (elicitation) of human utilities:
  - Compare a prize  $A$  to a standard lottery  $L_p$  between
    - ▶ “best possible prize”  $u_+$  with probability  $p$
    - ▶ “worst possible catastrophe”  $u_-$  with probability  $1 - p$
  - Adjust lottery probability  $p$  until indifference:  $A \sim L_p$
  - Resulting  $p$  is a utility in  $[0, 1]$



# Human Utilities (Revisited)

- Utilities map states to real numbers
- Standard approach to assessment (elicitation) of human utilities:
  - Compare a prize  $A$  to a standard lottery  $L_p$  between
    - ▶ “best possible prize”  $u_+$  with probability  $p$
    - ▶ “worst possible catastrophe”  $u_-$  with probability  $1 - p$
  - Adjust lottery probability  $p$  until indifference:  $A \sim L_p$
  - Resulting  $p$  is a utility in  $[0, 1]$




*Pay \$30*




# Human Utilities (Revisited)

- Utilities map states to real numbers
- Standard approach to assessment (elicitation) of human utilities:
  - Compare a prize  $A$  to a standard lottery  $L_p$  between
    - ▶ “best possible prize”  $u_+$  with probability  $p$
    - ▶ “worst possible catastrophe”  $u_-$  with probability  $1 - p$
  - Adjust lottery probability  $p$  until indifference:  $A \sim L_p$
  - Resulting  $p$  is a utility in  $[0, 1]$



Pay \$30




# Human Utilities (Revisited)

- Utilities map states to real numbers
- Standard approach to assessment (elicitation) of human utilities:
  - Compare a prize  $A$  to a standard lottery  $L_p$  between
    - ▶ “best possible prize”  $u_+$  with probability  $p$
    - ▶ “worst possible catastrophe”  $u_-$  with probability  $1 - p$
  - Adjust lottery probability  $p$  until indifference:  $A \sim L_p$
  - Resulting  $p$  is a utility in  $[0, 1]$




*Pay \$30*



# Human Utilities (Revisited)

- Utilities map states to real numbers
- Standard approach to assessment (elicitation) of human utilities:
  - Compare a prize  $A$  to a standard lottery  $L_p$  between
    - ▶ “best possible prize”  $u_+$  with probability  $p$
    - ▶ “worst possible catastrophe”  $u_-$  with probability  $1 - p$
  - Adjust lottery probability  $p$  until indifference:  $A \sim L_p$
  - Resulting  $p$  is a utility in  $[0, 1]$





*Pay \$30*

*No change*

*Instant death*



# Human Utilities (Revisited)

- Utilities map states to real numbers
- Standard approach to assessment (elicitation) of human utilities:
  - Compare a prize  $A$  to a standard lottery  $L_p$  between
    - ▶ “best possible prize”  $u_+$  with probability  $p$
    - ▶ “worst possible catastrophe”  $u_-$  with probability  $1 - p$
  - Adjust lottery probability  $p$  until indifference:  $A \sim L_p$
  - Resulting  $p$  is a utility in  $[0, 1]$



# Human Utilities (Revisited)

- Utilities map states to real numbers
- Standard approach to assessment (elicitation) of human utilities:
  - Compare a prize  $A$  to a standard lottery  $L_p$  between
    - ▶ “best possible prize”  $u_+$  with probability  $p$
    - ▶ “worst possible catastrophe”  $u_-$  with probability  $1 - p$
  - Adjust lottery probability  $p$  until indifference:  $A \sim L_p$
  - Resulting  $p$  is a utility in  $[0, 1]$



# Money

- Money does not behave as a utility function, but we can talk about the utility of having money (or being in debt)



# Money

- Money does not behave as a utility function, but we can talk about the utility of having money (or being in debt)
- Given a lottery  $L = [p, \$X; (1 - p), \$Y]$ 
  - The **expected monetary value**  $EMV(L)$  is  $p \times X + (1 - p) \times Y$
  - $U(L) = p \times U(\$X) + (1 - p) \times U(\$Y)$
  - Typically,  $U(L) < EMV(L)$



# Money

- Money does not behave as a utility function, but we can talk about the utility of having money (or being in debt)
- Given a lottery  $L = [p, \$X; (1 - p), \$Y]$ 
  - The **expected monetary value**  $EMV(L)$  is
$$p \times X + (1 - p) \times Y$$
  - $U(L) = p \times U(\$X) + (1 - p) \times U(\$Y)$
  - Typically,  $U(L) < U(EMV(L))$
- People are **risk-averse**






# Money

- Money does not behave as a utility function, but we can talk about the utility of having money (or being in debt)
- Given a lottery  $L = [p, \$X; (1 - p), \$Y]$ 
  - The **expected monetary value**  $EMV(L)$  is
$$p \times X + (1 - p) \times Y$$
  - $U(L) = p \times U(\$X) + (1 - p) \times U(\$Y)$
  - Typically,  $U(L) < U(EMV(L))$
- People are **risk-averse**
- When deep in debt, people are **risk-prone**



# Money

- Money does not behave as a utility function, but we can talk about the utility of having money (or being in debt)
- Given a lottery  $L = [p, \$X; (1 - p), \$Y]$ 
  - The **expected monetary value**  $EMV(L)$  is  $p \times X + (1 - p) \times Y$
  - $U(L) = p \times U(\$X) + (1 - p) \times U(\$Y)$
  - Typically,  $U(L) < U(EMV(L))$
- People are **risk-averse**
- When deep in debt, people are **risk-prone**



# Suggested Reading

- Russell & Norvig: Chapter 5.2-5.5, 16.1-16.3