Artificial Intelligence CSE 4617

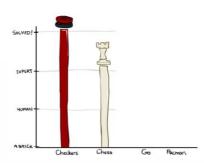
Ahnaf Munir Assistant Professor Islamic University of Technology

Game Playing (State-of-the-Art)

- Checkers
 - 1950: First computer player
 - 1994: First computer champion
 - Chinook ended 40-year-reign of human champion Marion Tinsley using complete 8-piece endgame
 - 2007: Solved
 - If both players play optimally, you can at least force a draw

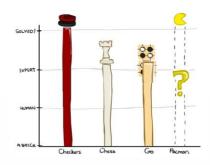
Chess

- 1997: Deep Blue defeats human champion Gary Kasparove in a six-game match
 - Examined 200M positions per second
 - Used very sophisticated evaluation function
 - Undisclosed methods for searching up to 40 ply
- Current programs are even better, if less historic



Game Playing (State-of-the-Art)

- Go
 - Human champions are now starting to be challenged by machines
 - Branching Factor b > 300
 - Classic programs → Pattern knowledge bases
 - Recent programs → Monte Carlo (randomized) expansion methods
 - 2016: Alpha GO defeats human champiom
 - Uses Monte Carlo Tree Search
 - Learned evaluation function: Odd-Even Function
- Pacman



Adversarial Games



Types of Games

- Many different kinds of games!
- Criteria/Axes:
 - Deterministic or stochastic?
 - e.g., Chess vs Monopoly
 - One, two, or more players?
 - e.g., Solitaire vs Checkers vs D&D, etc.
 - Zero sum?
 - e.g., Football vs Nuclear war
 - Perfect information?
 - e.g., Tic-Tac-Toe vs Poker
- Want algorithms for calculating a strategy (policy) which recommends a move from each state



Deterministic Games

- Many possible formalizations, one is:
 - States: S (start at s_0)
 - Players: $P = \{1...N\}$ (usually take turns)
 - Actions: *A* (may depend on player/state)
 - Transition Function: $S \times A \rightarrow S$
 - Terminal Test: $S \rightarrow \{t, f\}$
 - Terminal Utilities: $S \times P \rightarrow R$
- Solution for a player is a **policy**: $S \rightarrow A$

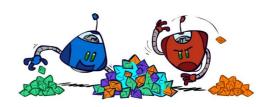


Zero-Sum Games



Zero-Sum Games

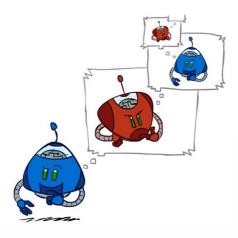
- Agents have opposite utilities (values on outcomes)
- Lets us think of a single value that one maximizes and the other minimizes
- Adversarial, pure competition



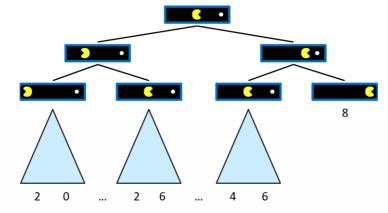
General Games

- Agents have independent utilities (values on outcomes)
- Cooperation, indifference, competition, and more are all possible
- More later on non-zero-sum games

Adversarial Search

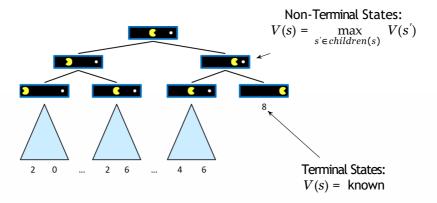


Single-Agent Trees

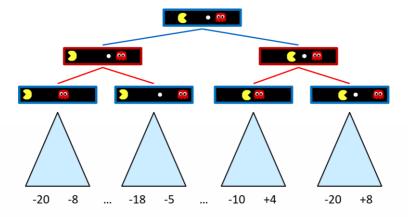


Value of a State

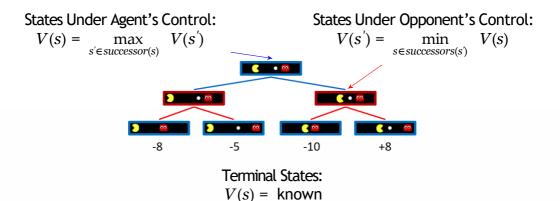
■ The best achievable outcome (utility) from that state



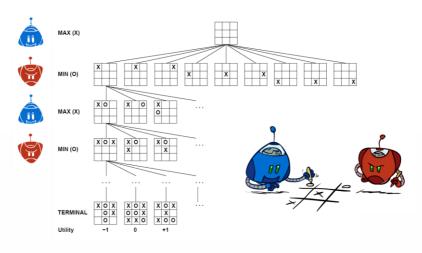
Adversarial Game Trees



Minimax Values

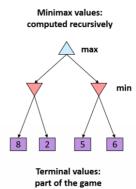


Tic-Tac-Toe Game Tree



Adversarial Search (Minimax)

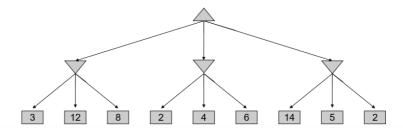
- Deterministic, zero-sum games:
 - Tic-tac-toe, chess, checkers
 - One player maximizes result
 - The other minimizes result
- Minimax search:
 - A state-space search tree
 - Players alternate turns
 - Compute each node's minimax value: the best achievable utility against a rational (optimal) adversary



Minimax Implementation (Dispatch)

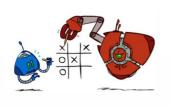
```
def value(state):
  if the state is a terminal state: return the state's utility
  if the next agent is MAX: return max-value(state)
  if the next agent is MIN: return min-value(state)
def max-value(state):
                                             def min-value(state):
    initialize v = -\infty
                                                initialize v = +\infty
   for each successor of state:
                                                for each successor of state:
       v = \max(v, \text{value}(successor))
                                                   v = \min(v, \text{value}(successor))
    return v
                                                return v
                           V(s')
   V(s) = \max
                                                                          V(s')
```

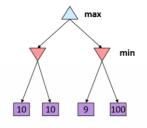
Minimax Example

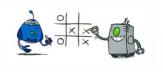


Minimax Properties

Optimal against a perfect player. Otherwise?



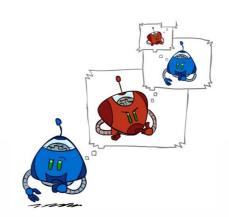




Video: min, exp

Minimax Efficiency

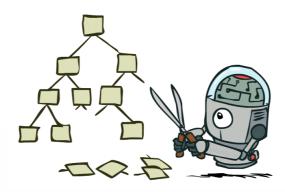
- How efficient is minimax?
 - Just like (exhaustive) DFS
 - Time: $O(b^m)$
 - Space: *O*(*bm*)
- **Example:** For chess, $b \approx 35$, $m \approx 100$
 - Exact solution is completely infeasible
 - But, do we need to explore the whole tree?



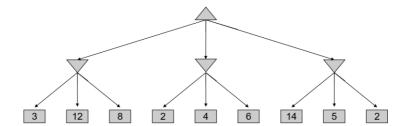
Resource Limits



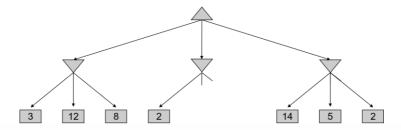
Game Tree Pruning



Minimax Example (Revisited)

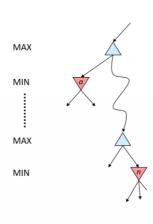


Minimax Pruning



Alpha-Beta Pruning

- General configuration (MIN version)
 - Computing the MIN-VALUE at some node n
 - Looping over n's children
 - n's estimate of the children's min is dropping
 - Who cares about n's value? MAX
 Let a be the best value that MAX can
 - get at any choice point along the current path from the root
 - If n becomes worse than a, MAX will avoid it, so we can stop considering n's other children (it's already bad enough that it won't be played)
- MAX version is symmetric



Alpha-Beta Implementation

 α : MAX's best option on path to root β : MIN's best option on path to root

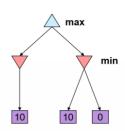
```
def max-value(state, \alpha, \beta):
    initialize v = -\infty
    for each successor of state:

v = \max(v, value(successor, \alpha, \beta))
    if v \ge \beta: return v
    \alpha = \max(\alpha, v)
    return v

def min-value(state, \alpha, \beta):
    initialize v = +\infty
    for each successor of state:
    v = \min(v, value(successor, \alpha, \beta))
    if v \le \alpha: return v
    \beta = \min(\beta, v)
    return v
```

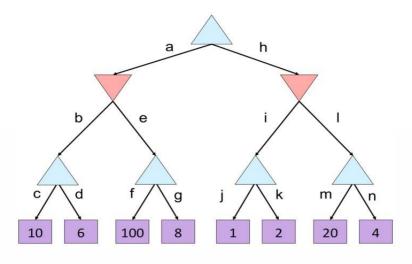
Alpha-Beta Pruning Properties

- The pruning has no effect on minimax value computed for the root!
- Values of intermediate nodes might be wrong
 - Important: children of the root may have the wrong value
 - The most naïve version won't let you do action selection
 - Solution 1: Prune only on inequality
 - Solution 2: Keep track of which one was first
- Good child ordering improves effectiveness
- With "perfect ordering":
 - Time complexity drops to $O(b^{m/2})$
 - Doubles solvable depth!
 - Full search of, e.g. chess, is still hopeless...
- This is a simple example of metareasoning (computing about what to compute)



Alpha-Beta Quiz е 50 10

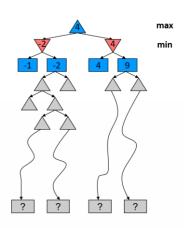
Alpha-Beta Quiz



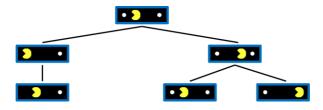
Resource Limits

- Problem: In realistic games, cannot search to leaves!
- Solution: Depth-limited search
 - Search only to a limited depth in the tree
 - Replace terminal utilities with an evaluation function for non-terminal positions
- Example:
 - Suppose we have 100 seconds, can explore 10K nodes/sec
 - Can check 1M nodes per move
 - $\alpha \beta$ reaches about depth 8 decent chess program
- Guarantee of optimal play is gone
- More plies makes a BIG difference
- Use iterative deepening for an anytime algorithm

Video: demo-thrashing



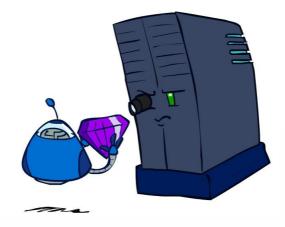
Why Pacman Starves



- A danger of replanning agents!
 - He knows his score will go up by eating the dot now (west, east)
 - He knows his score will go up just as much by eating the dot later (east, west)
 - There are no point-scoring opportunities after eating the dot (within the horizon, two here)
 - Therefore, waiting seems just as good as eating: he may go east, then back west in the next round of replanning!

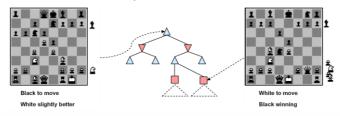
Video: demo-thrashing-fixed

Evaluation Functions

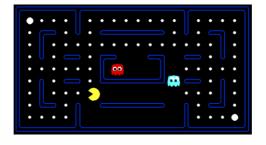


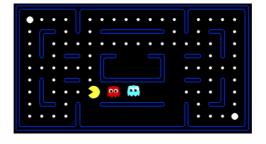
Evaluation Functions

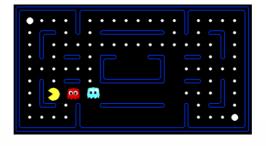
Used to score non-terminals in depth-limited search

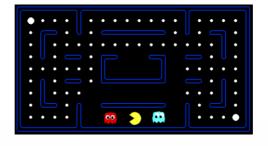


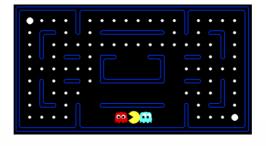
- Ideal function: returns the actual minimax value of the position
- In practice: typically weighted linear sum of features: $Eval(s) = w_1f_1(s) + w_2f_2(s) + \cdots + w_nf_n(s)$
- e.g.: $f_1(s)$ = (# of white queens # of black queens), etc.











Depth Matters

- Evaluation functions are always imperfect
- The deeper in the tree the evaluation function is buried, the less the quality of the evaluation function matters
- An important example of the tradeoff between complexity of features and complexity of computation





Video: depth-limited-1, depth-limited-10

Suggested Reading

■ Russell & Norvig: Chapter 5.2-5.5