

Artificial Intelligence

CSE 4617

Ahnaf Munir
Assistant Professor
Islamic University of Technology


Game Playing (State-of-the-Art)

■ Checkers

- 1950: First computer player
- 1994: First computer champion
 - ▶ Chinook ended 40-year-reign of human champion Marion Tinsley using complete 8-piece endgame
- 2007: Solved
 - ▶ If both players play optimally, you can at least force a draw

■ Chess

- 1997: Deep Blue defeats human champion Gary Kasparov in a six-game match
 - ▶ Examined 200M positions per second
 - ▶ Used very sophisticated evaluation function
 - ▶ Undisclosed methods for searching up to 40 ply
- Current programs are even better, if less historic




Game Playing (State-of-the-Art)

■ Go

- Human champions are now starting to be challenged by machines
 - ▶ Branching Factor $b > 300$
 - ▶ Classic programs → Pattern knowledge bases
 - ▶ Recent programs → Monte Carlo (randomized) expansion methods
- 2016: Alpha Go defeats human champion
 - ▶ Uses Monte Carlo Tree Search
 - ▶ Learned evaluation function: Odd-Even Function

■ Pacman



Adversarial Games




Types of Games

- Many different kinds of games!
- Criteria/Axes:
 - Deterministic or stochastic?
 - ▶ e.g., Chess vs Monopoly
 - One, two, or more players?
 - ▶ e.g., Solitaire vs Checkers vs D&D, etc.
 - Zero sum?
 - ▶ e.g., Football vs Nuclear war
 - Perfect information?
 - ▶ e.g., Tic-Tac-Toe vs Poker
- Want algorithms for calculating a **strategy (policy)** which recommends a move from each state



Deterministic Games

- Many possible formalizations, one is:
 - States: S (start at s_0)
 - Players: $P = \{1 \dots N\}$ (usually take turns)
 - Actions: A (may depend on player/state)
 - Transition Function: $S \times A \rightarrow S$
 - Terminal Test: $S \rightarrow \{t, f\}$
 - Terminal Utilities: $S \times P \rightarrow R$
- Solution for a player is a **policy**: $S \rightarrow A$



Zero-Sum Games



■ Zero-Sum Games


- Agents have opposite utilities (values on outcomes)
- Lets us think of a single value that one maximizes and the other minimizes
- Adversarial, pure competition




■ General Games

- Agents have independent utilities (values on outcomes)
- Cooperation, indifference, competition, and more are all possible
- More later on non-zero-sum games

Adversarial Search




Single-Agent Trees




Value of a State

- The best achievable outcome (utility) from that state



Adversarial Game Trees




Minimax Values

States Under Agent's Control:

$$V(s) = \max_{s' \in \text{successor}(s)} V(s')$$

States Under Opponent's Control:


$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$



Terminal States:


$$V(s) = \text{known}$$

Tic-Tac-Toe Game Tree



Adversarial Search (Minimax)

- Deterministic, zero-sum games:
 - Tic-tac-toe, chess, checkers
 - One player maximizes result
 - The other minimizes result
- Minimax search:
 - A state-space search tree
 - Players alternate turns
 - Compute each node's minimax value: the best achievable utility against a rational (optimal) adversary



Minimax Implementation (Dispatch)

def value(state):

if the state is a terminal state: return the state's utility

if the next agent is MAX: return max-value(state)

if the next agent is MIN: return min-value(state)

def max-value(state):

initialize $v = -\infty$

for each successor of *state*:

$v = \max(v, \text{value}(\text{successor}))$

return v

def min-value(state):

initialize $v = +\infty$

for each successor of *state*:


$v = \min(v, \text{value}(\text{successor}))$

return v

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$




$$V(s) = \min_{s' \in \text{successors}(s)} V(s')$$

Minimax Example



Minimax Properties


Optimal against a perfect player. Otherwise?



Video: [min](#), [exp](#)

Minimax Efficiency


- How efficient is minimax?
 - Just like (exhaustive) DFS
 - Time: $O(b^m)$
 - Space: $O(bm)$
- Example: For chess, $b \approx 35$, $m \approx 100$
 - Exact solution is completely infeasible
 - But, do we need to explore the whole tree?




Resource Limits




Game Tree Pruning



Minimax Example (Revisited)




Minimax Pruning



Alpha-Beta Pruning

- General configuration (MIN version)
 - Computing the MIN-VALUE at some node n
 - Looping over n 's children
 - n 's estimate of the children's min is dropping
 - Who cares about n 's value? MAX
 - Let a be the best value that MAX can get at any choice point along the current path from the root
 - If n becomes worse than a , MAX will avoid it, so we can stop considering n 's other children (it's already bad enough that it won't be played)
- MAX version is symmetric



Alpha-Beta Implementation

α : MAX's best option on path to root

β : MIN's best option on path to root

def max-value(*state*, α , β):

 initialize $v = -\infty$

 for each successor of *state*:

$v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$

 if $v \geq \beta$: return v

$\alpha = \max(\alpha, v)$

 return v

def min-value(*state*, α , β):

 initialize $v = +\infty$

 for each successor of *state*:

$v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$


 if $v \leq \alpha$: return v

$\beta = \min(\beta, v)$


 return v

Alpha-Beta Pruning Properties


- The pruning has **no effect** on minimax value computed for the root!
- Values of intermediate nodes might be wrong
 - Important: children of the root may have the wrong value
 - The most naïve version won't let you do action selection
 - ▶ Solution 1: Prune only on inequality
 - ▶ Solution 2: Keep track of which one was first
- Good child ordering improves effectiveness
- With “perfect ordering”:
 - Time complexity drops to $O(b^{m/2})$
 - Doubles solvable depth!
 - Full search of, e.g. chess, is still hopeless...
- This is a simple example of **metareasoning** (computing about what to compute)



Alpha-Beta Quiz




Alpha-Beta Quiz




Resource Limits

- Problem: In realistic games, cannot search to leaves!
- Solution: Depth-limited search
 - Search only to a limited depth in the tree
 - Replace terminal utilities with an evaluation function for non-terminal positions
- Example:
 - Suppose we have 100 seconds, can explore 10K nodes/sec
 - Can check 1M nodes per move
 - $\alpha - \beta$ reaches about depth 8 - decent chess program
- Guarantee of optimal play is gone
- More plies makes a BIG difference
- Use iterative deepening for an anytime algorithm

Video: [demo-thrashing](#)



Why Pacman Starves




■ A danger of replanning agents!

- He knows his score will go up by eating the dot now (west, east)
- He knows his score will go up just as much by eating the dot later (east, west)
- There are no point-scoring opportunities after eating the dot (within the horizon, two here)
- Therefore, waiting seems just as good as eating: he may go east, then back west in the next round of replanning!


Video: [demo-thrashing-fixed](#)

Evaluation Functions




Evaluation Functions

- Used to score non-terminals in depth-limited search




- Ideal function: returns the actual minimax value of the position
- In practice: typically weighted linear sum of features:
$$Eval(s) = w_1f_1(s) + w_2f_2(s) + \dots + w_nf_n(s)$$
- e.g.: $f_1(s) = (\# \text{ of white queens} - \# \text{ of black queens}), \text{ etc.}$

Evaluation Function for Pacman




Video: [smart-ghosts](#), [smart-ghosts-zoomed](#)

Evaluation Function for Pacman




Video: [smart-ghosts](#), [smart-ghosts-zoomed](#)

Evaluation Function for Pacman




Video: [smart-ghosts](#), [smart-ghosts-zoomed](#)

Evaluation Function for Pacman



Video: [smart-ghosts](#), [smart-ghosts-zoomed](#)

Evaluation Function for Pacman



Video: [smart-ghosts](#), [smart-ghosts-zoomed](#)

Depth Matters

- Evaluation functions are always imperfect
- The deeper in the tree the evaluation function is buried, the less the quality of the evaluation function matters
- An important example of the tradeoff between complexity of features and complexity of computation



Video: [depth-limited-2](#), [depth-limited-10](#)

Suggested Reading

- Russell & Norvig: Chapter 5.2-5.5