


Database Denormalization

SWE 4601 Software Design and Architecture



Normalization & Denormalization

- **Normalization** is a technique of eliminating the redundant data from the database.
- **Denormalization** is the inverse process of normalization where the redundancy is added to the data to improve the performance of the specific application.
- By updating, inserting, and deleting records through foreign key relationships, normalization reduces instances of missed or orphaned records in your database.
- **Number of tables** **Increases in normalization**
 Decreases in denormalization
- While normalized data is optimized for entity level transactions, denormalized data is optimized for answering business questions and driving decision making.

Pros of Normalization

- Complies to ACID property
 - Atomicity
 - Consistency
 - Isolation
 - Durability
- Updates run quickly
- Inserts run quickly
- Clear understanding of data
- No redundancy

So, why denormalization?



Normalized DB require **join**

A lot of join

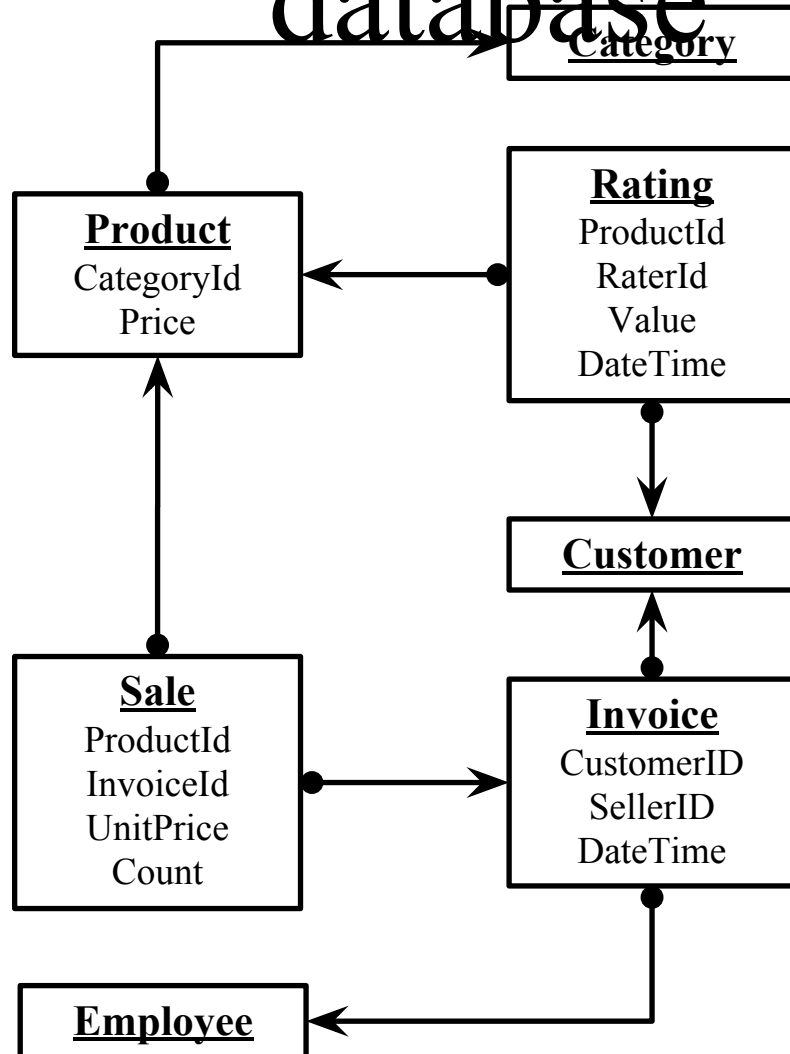
May be **crazy** lot of join

So What?

¯_(ツ)_/¯

Join is **expensive**

The kid's shop database



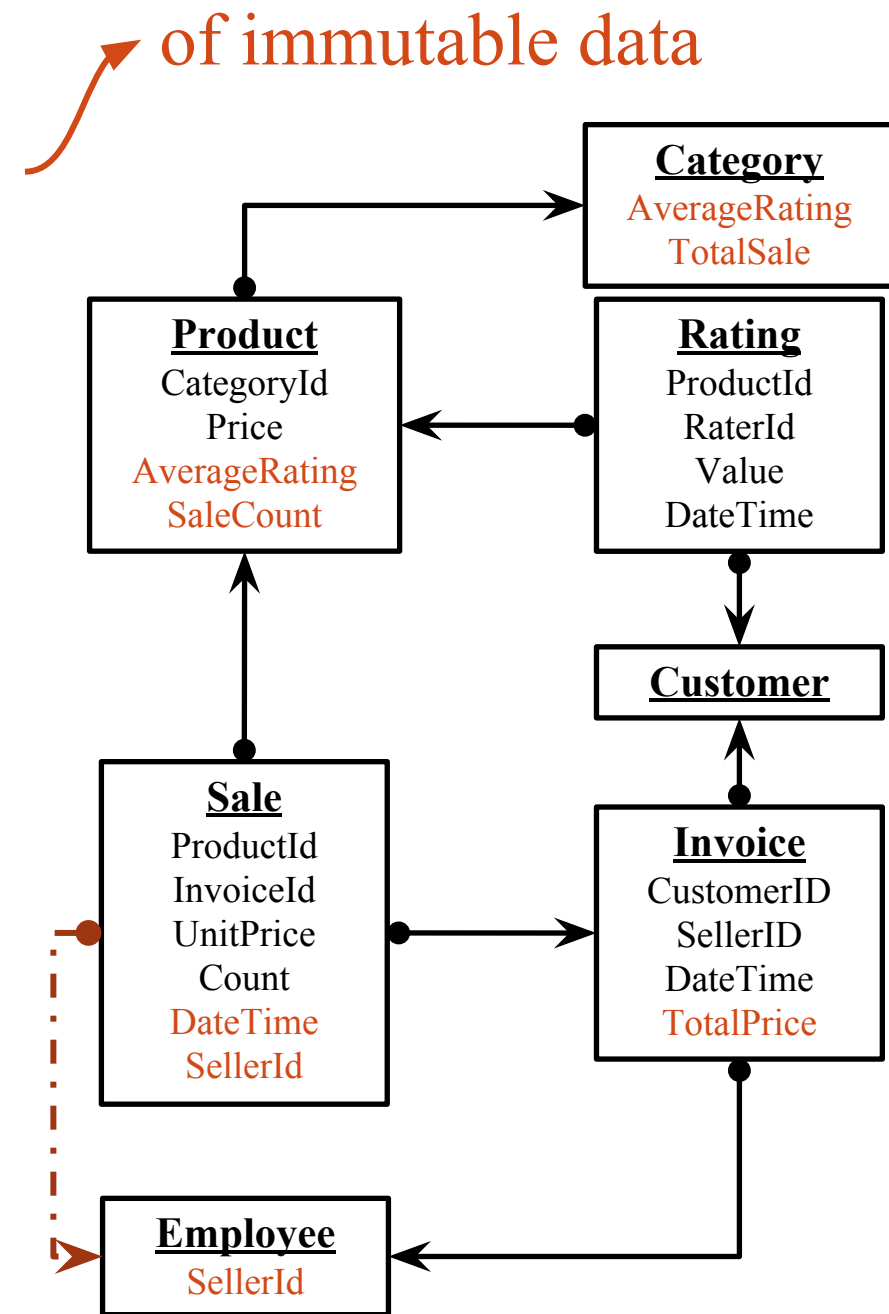
Which tables need to be joined for the following use cases?

1. A customer wants to see top rated products
2. A sales manager wants to find products with most number of sells
3. A sales manager wants to find products with most sale amount last month
4. A sales manager wants to award the top 3 salespersons of the last month
5. A marketing manager wants to award discount coupons to customers who purchased more than 10,000 Taka last month
6. A customer wants to see popular categories
7. Top management want to see a chart showing total sale of top 5 categories in a May 2020
8. A sales manager wants to see salespersons' sale per category

Speed up the queries by redundancy

1. A customer wants to see top rated products (Previously Product, Rating)
2. A sales manager wants to find products with most number of sales (Previously Product, Sale)
3. A sales manager wants to find products with most sale amount last month (Previously Product, Sale, Invoice)
4. A sales manager wants to award the top 3 sales persons of the last month (Previously Sale, Invoice, Employee)
5. A marketing manager wants to award discount coupons to customers who purchased more than 10000 Taka last month (Previously Customer, Invoice, Sale)
6. A customer wants to see popular categories (Previously Category, Product, Rating)
7. Top management want to see a chart showing total sale of top 5 categories in a May 2020 (Previously Sale, category, product, Invoice, rating)
8. A sales manager wants to see sales persons' sale per category (Previously Employee, Invoice, sale, Product, Category)

What if we keep a categoryId in Sale?

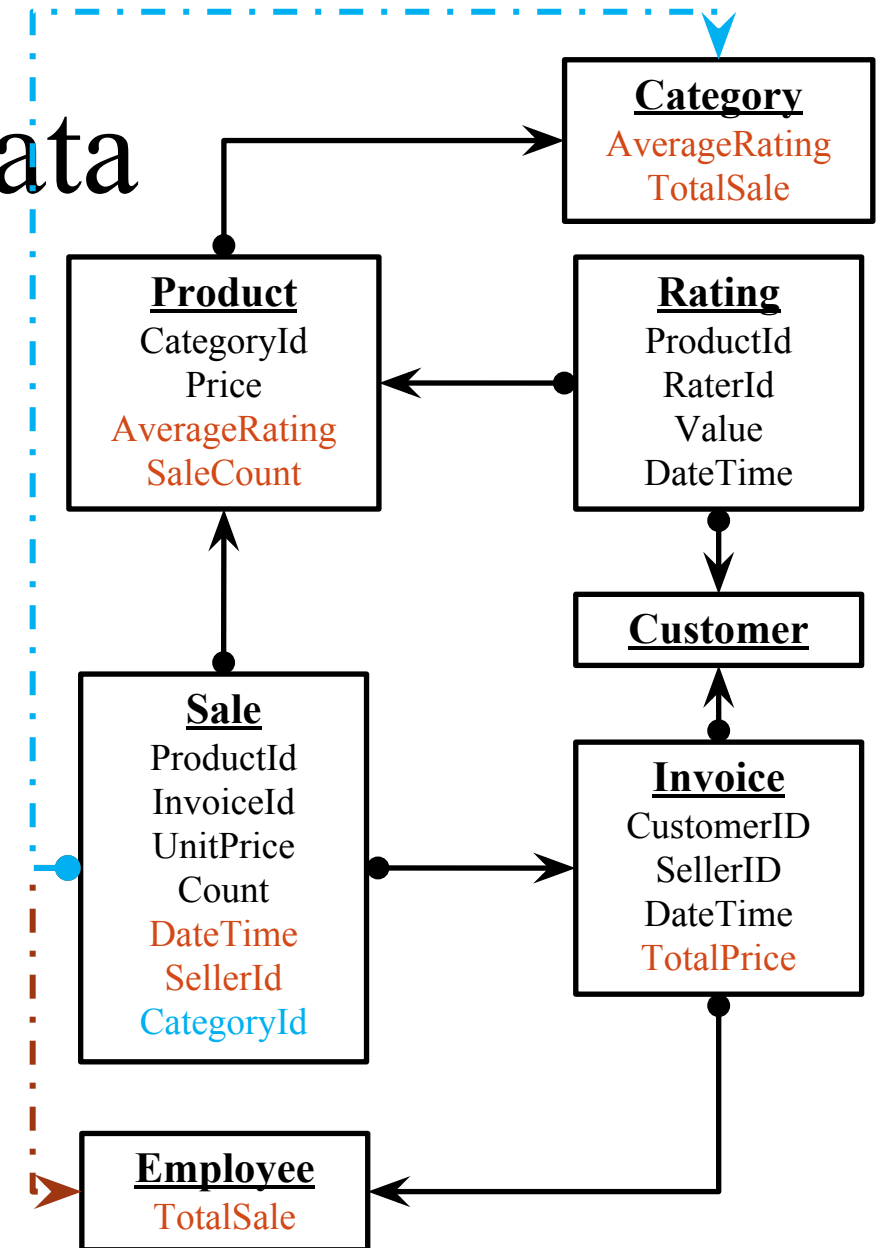


Redundancy of mutable data

- Wrong updates can be done from different parts of the code
 - Solution: Make sure update of one data is done by only one piece of code
- Updates can be slow
 - No big deal if update is not so frequent

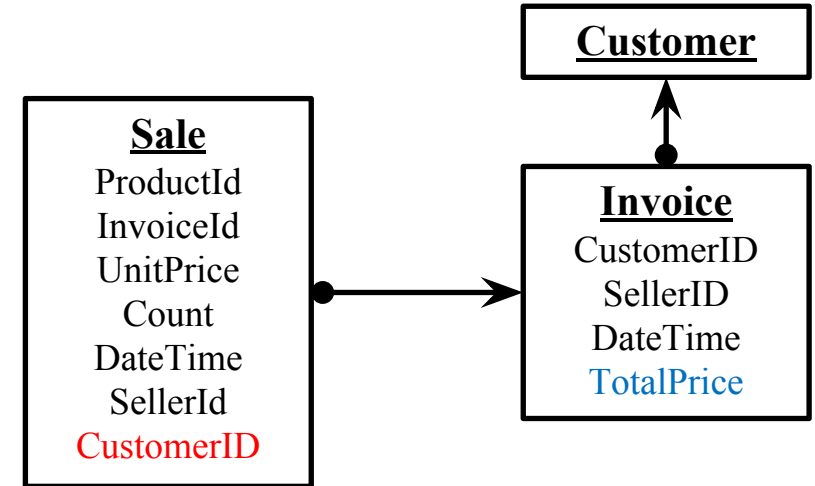
Lets solve now

7. Top management want to see a chart showing total sale of top 5 categories in a May 2020 (Previously Category, product, Sale)
8. A sales manager wants to see sales persons' sale per category (Previously Employee, Sale, Product, Category)



Choices need to be made

- A marketing manager wants to award discount coupons to customers who purchased more than 10,000 Taka last month (Customer, Invoice, Sale)



So, what did we gain by using redundancy?

- Faster query
- Easier query
- Meaningful query, examples
 - A customer wants to see top rated products
 - queries just product
 - A customer wants to see popular categories
 - queries just Category
 - Top management want to see a chart showing total sale of top 5 categories in a May 2020
 - queries Sale and Category
 - A sales manager wants to see sales persons' sale per category
 - queries Employee, Sale, Category